



Bayesian Statistics and Hierarchical Bayesian Modeling for Psychological Science

Lecture II

Lei Zhang

Social, Cognitive and Affective Neuroscience Unit (SCAN-Unit)
Department of Cognition, Emotion, and Methods in Psychology

https://github.com/lei-zhang/BayesCog_Wien

lei.zhang@univie.ac.at
lei-zhang.net
@lei_zhang_lz



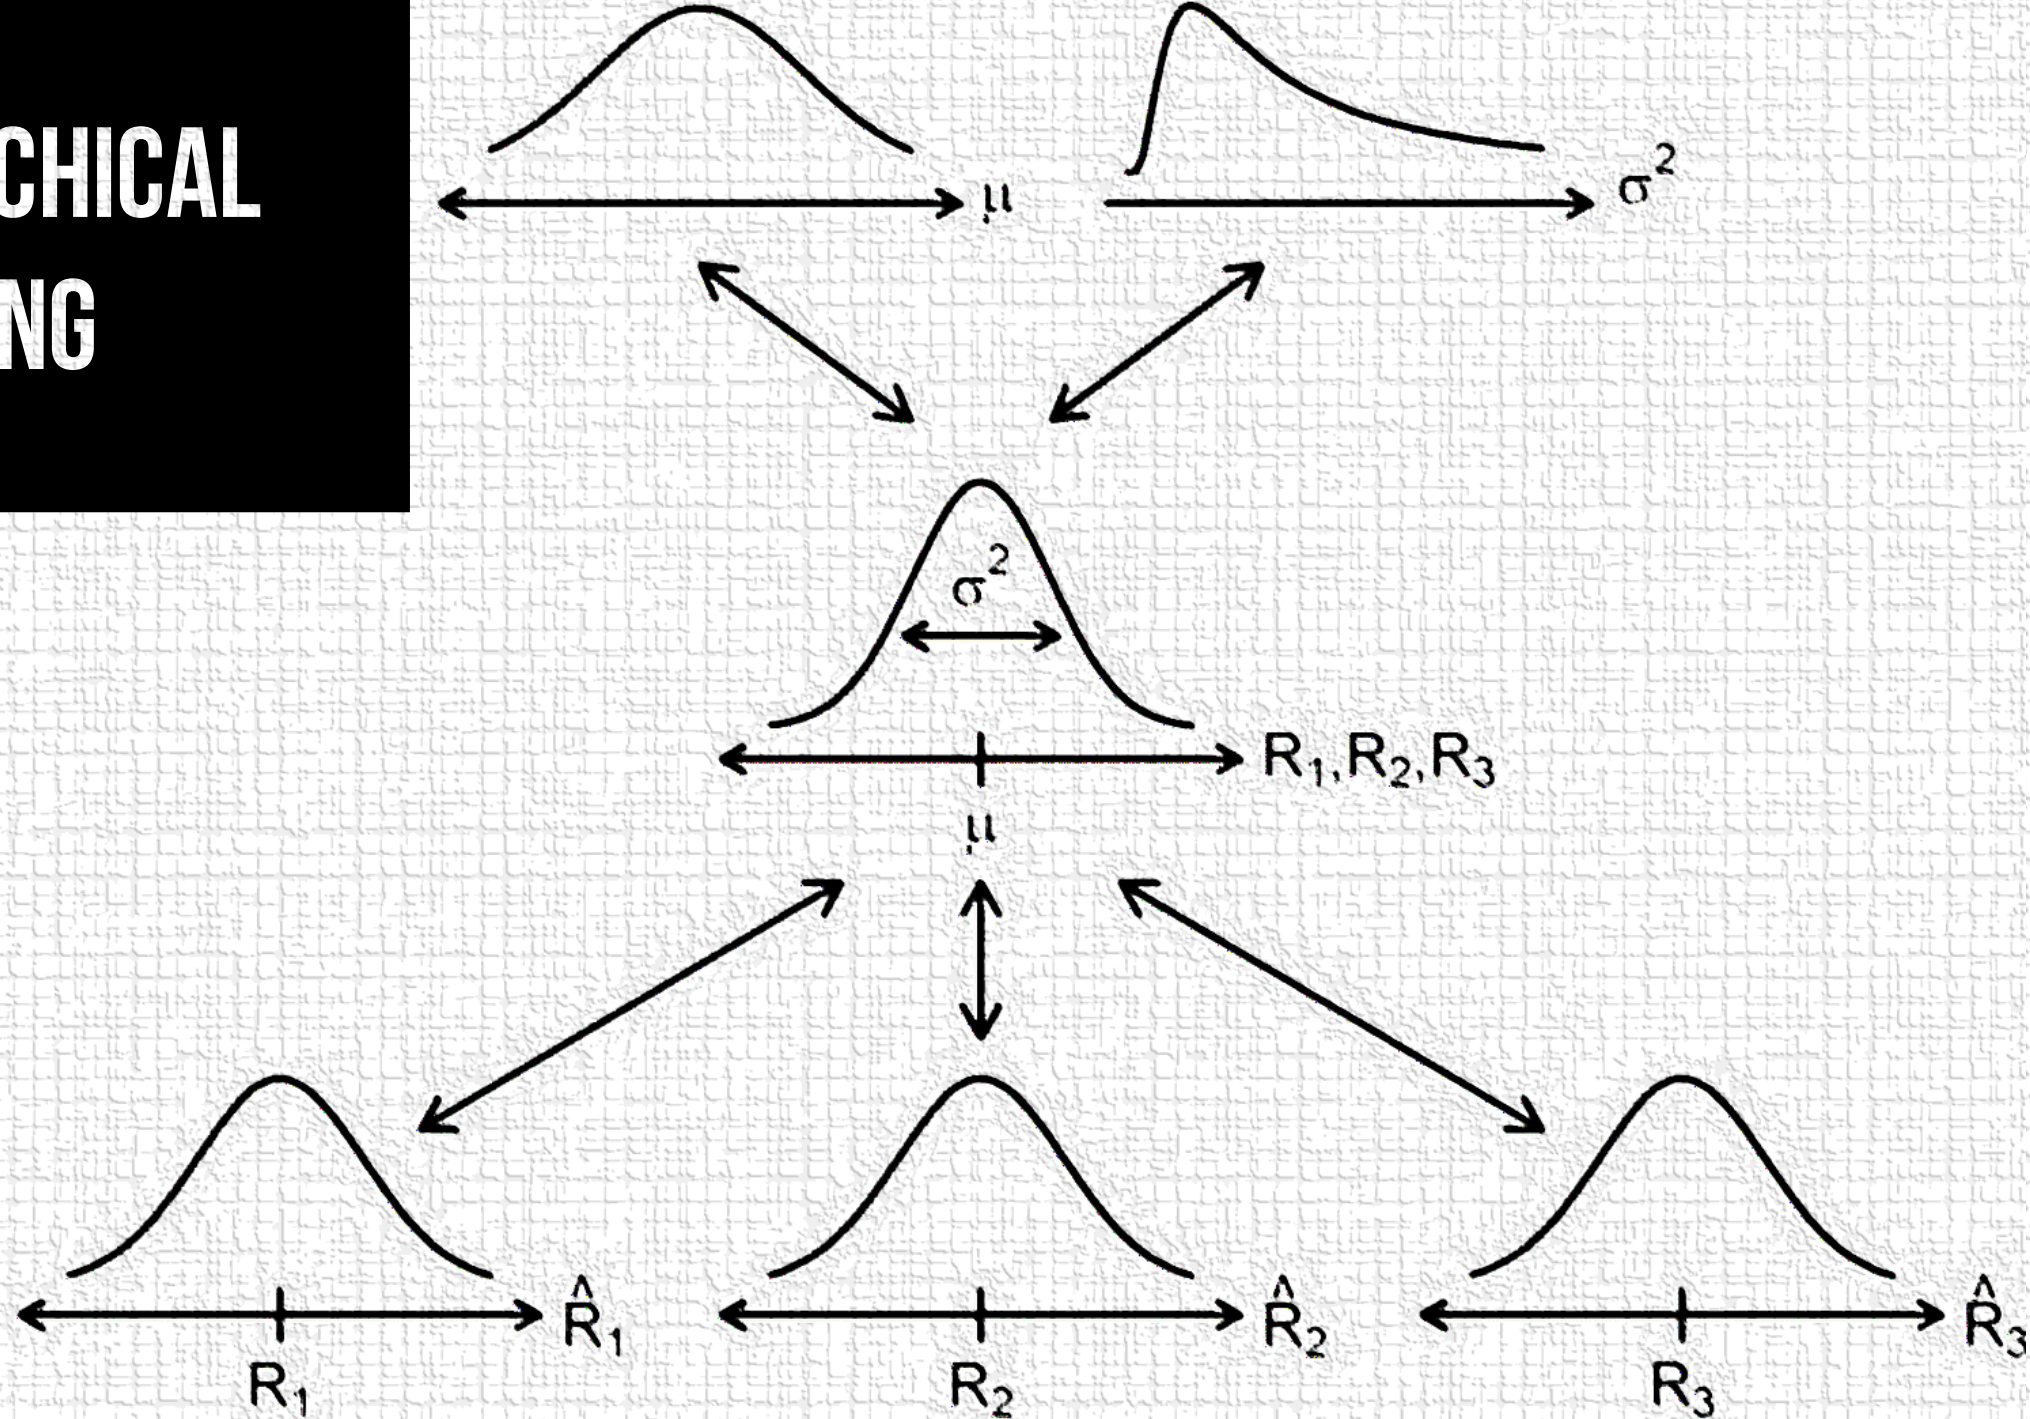
universität
wien

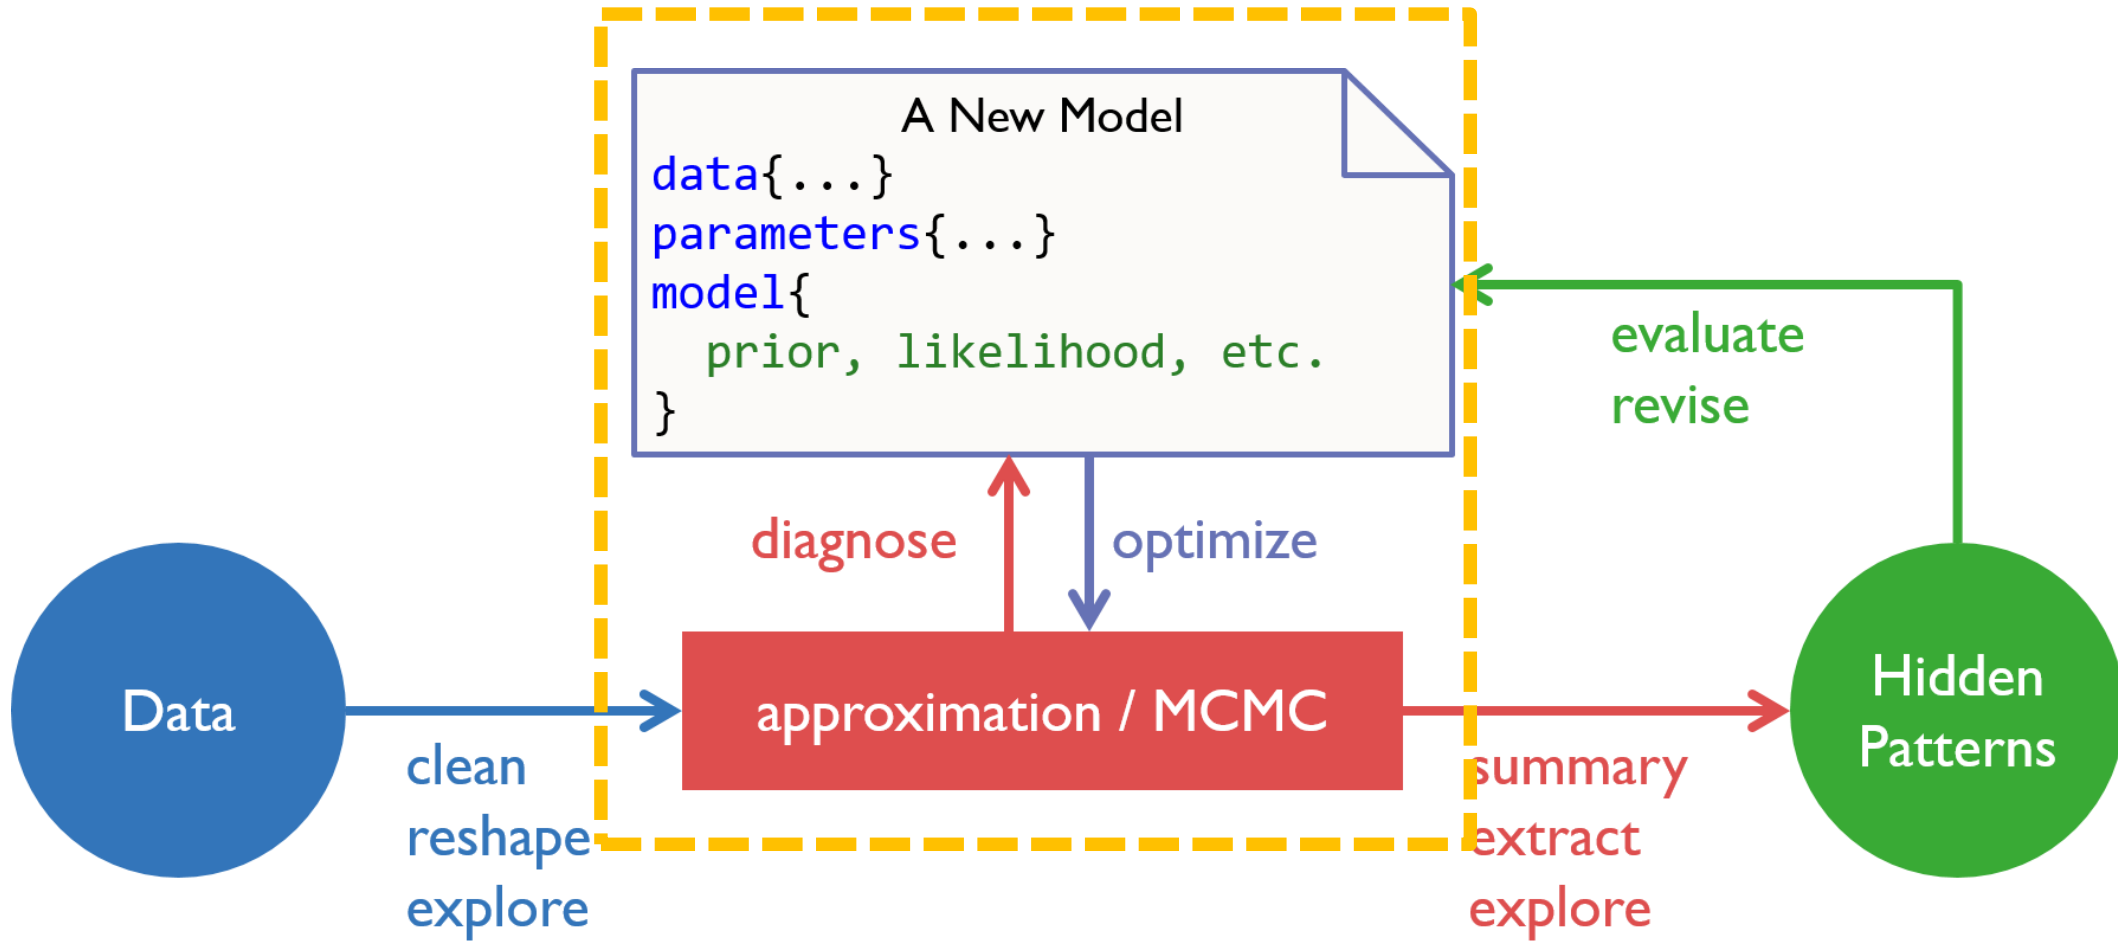
Fakultät für Psychologie



Bayesian warm-up?

HIERARCHICAL MODELING





Why Hierarchical Bayesian Cognitive Modeling?

cognitive model

statistics

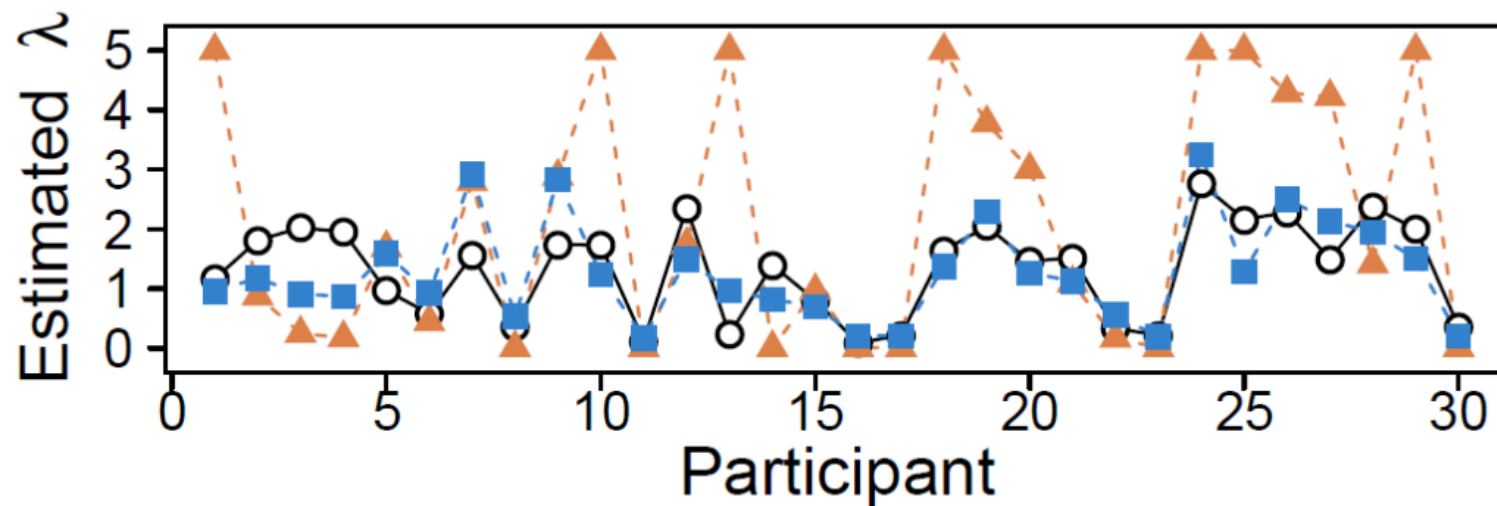
computing

Simulation study

Hierarchical Bayesian ■

Maximum likelihood ▲

Actual values ○



Why **Hierarchical** Bayesian Cognitive Modeling?

cognitive model

statistics

computing

Fixed effects

- all subjects are fitted with the **same set of parameters**
- worse model fit than “random effects”

Random effects

- each subject is fitted **independently of the others**
- best model fit for each subject
- parameter estimates can be noisy

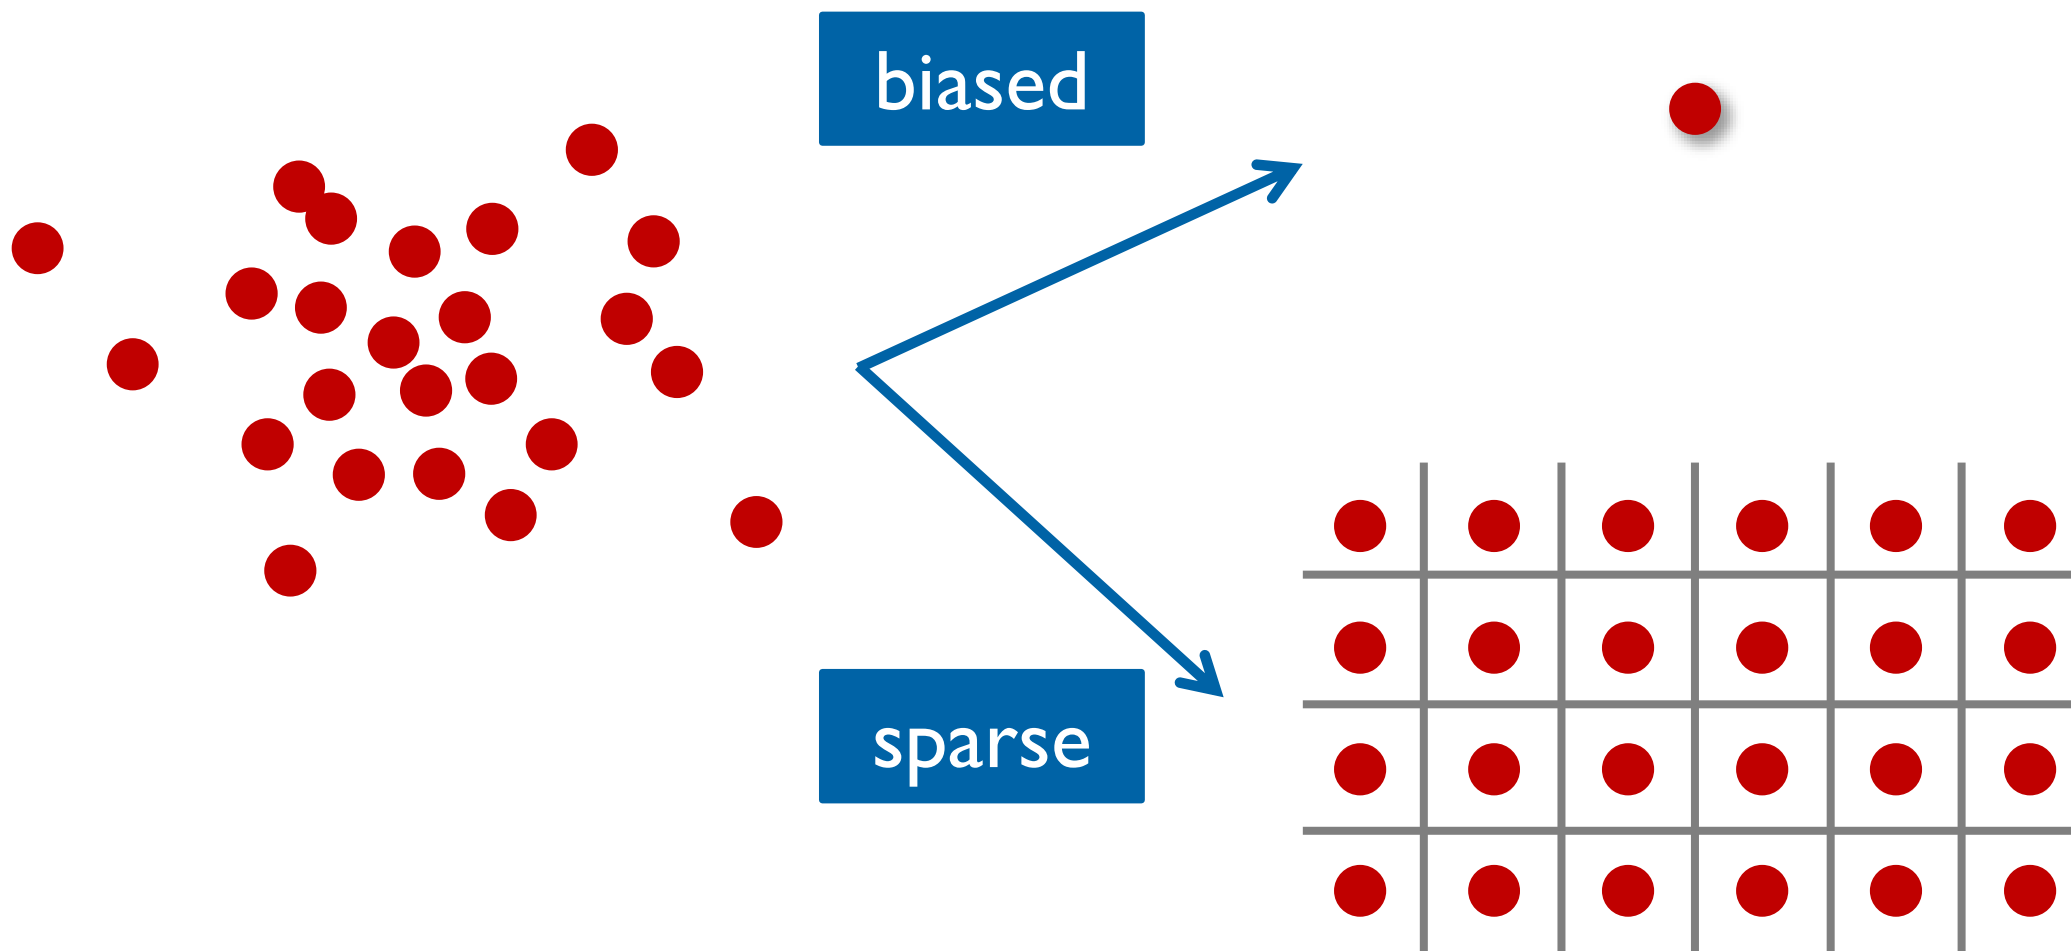
Adapted from Jan Gläscher's
workshop

Fitting Multiple Participants

cognitive model

statistics

computing

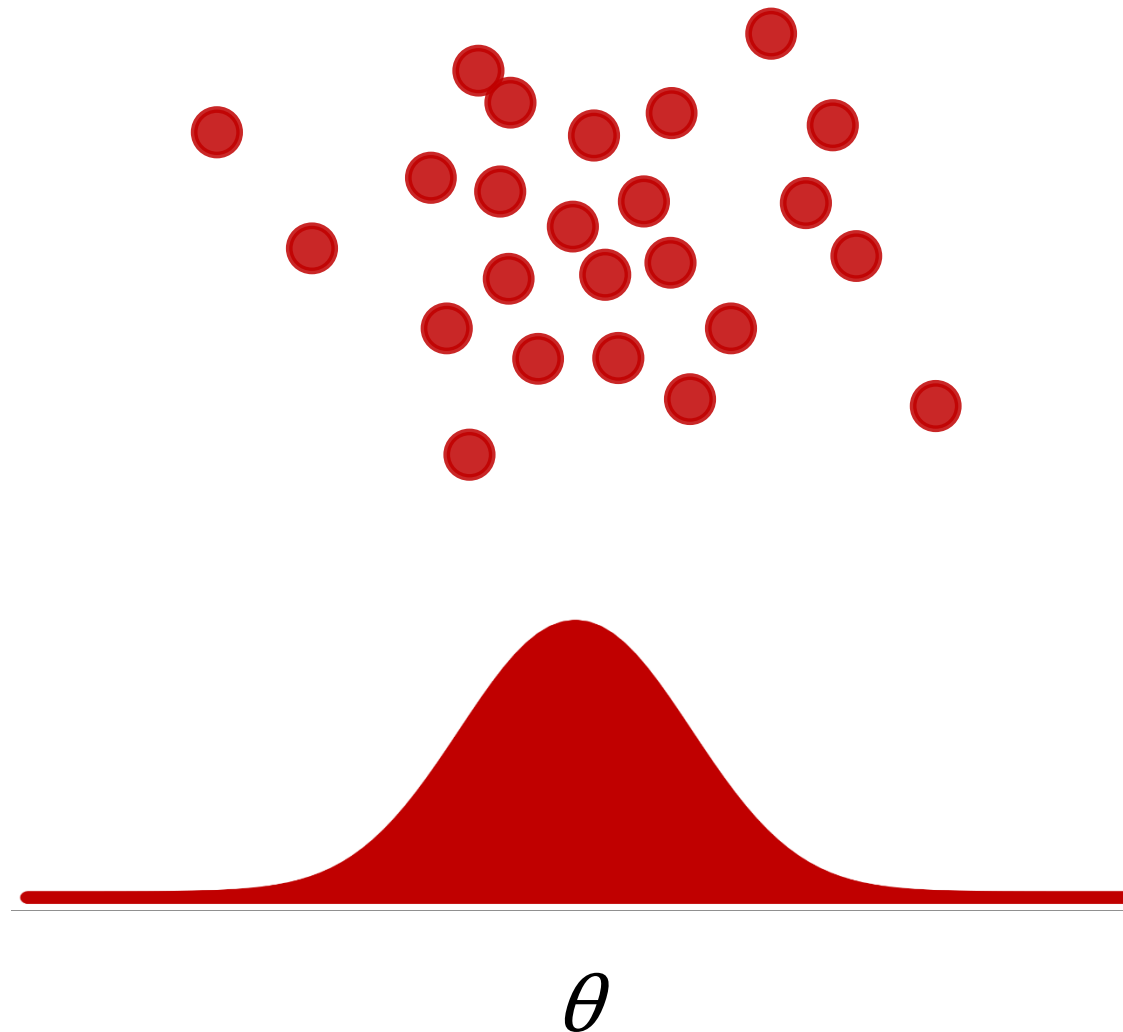


Fitting Multiple Participants

cognitive model

statistics

computing



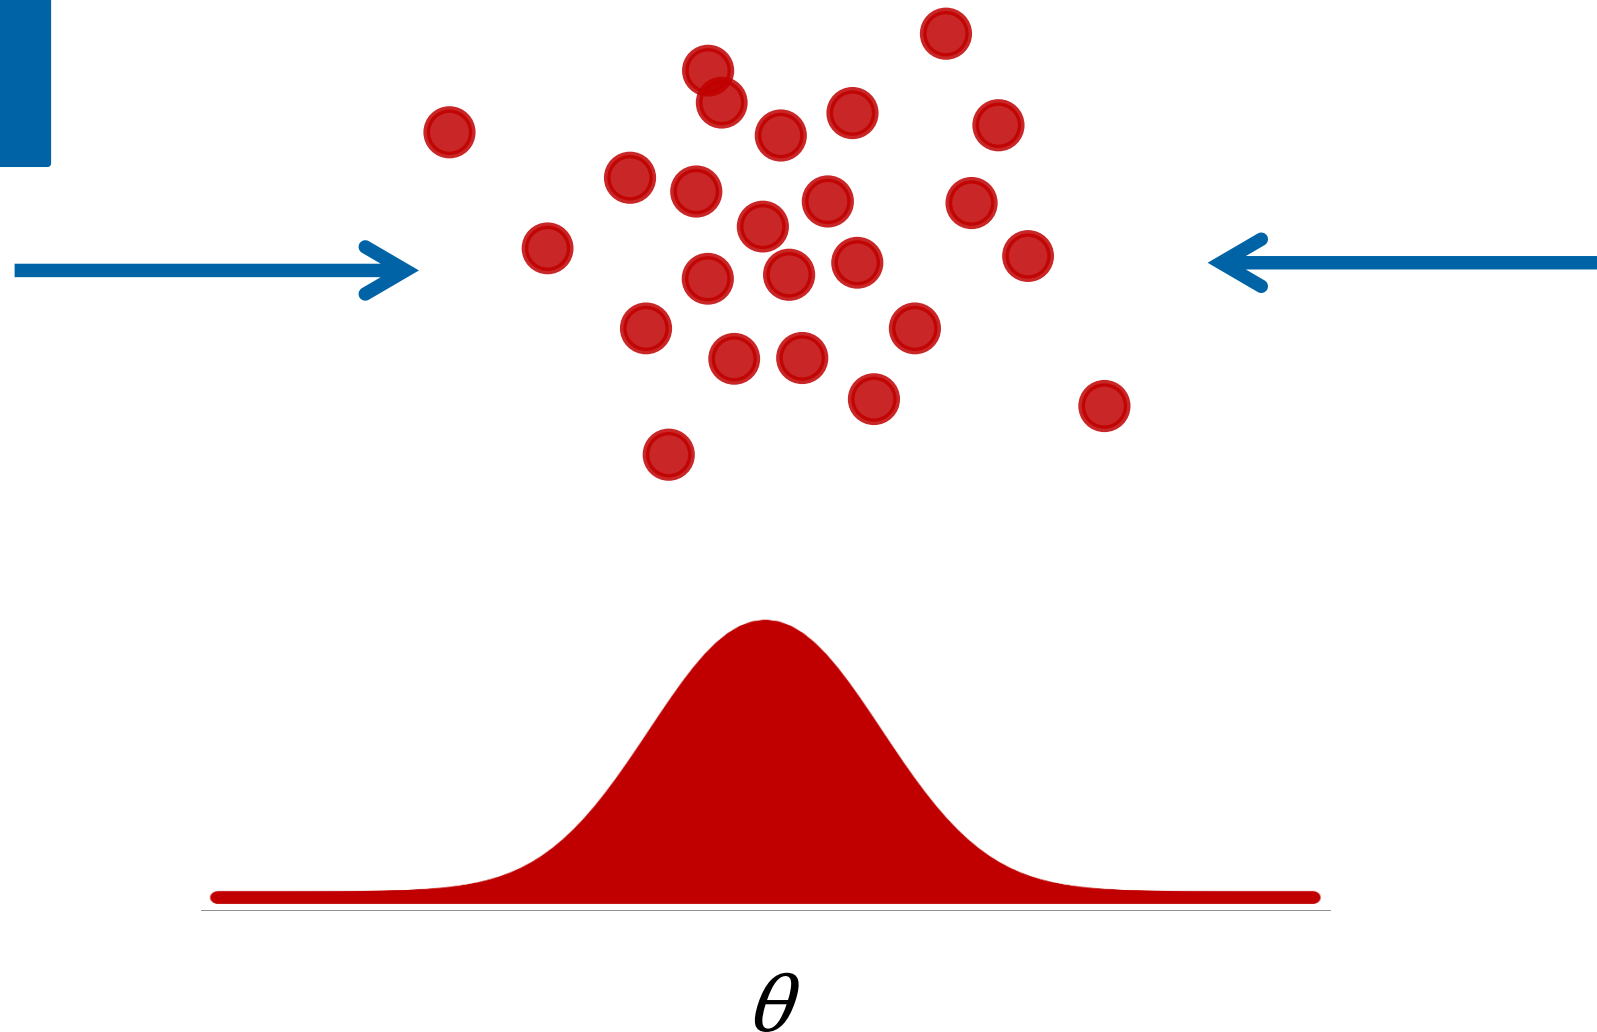
Fitting Multiple Participants

cognitive model

statistics

computing

shrinkage
effect

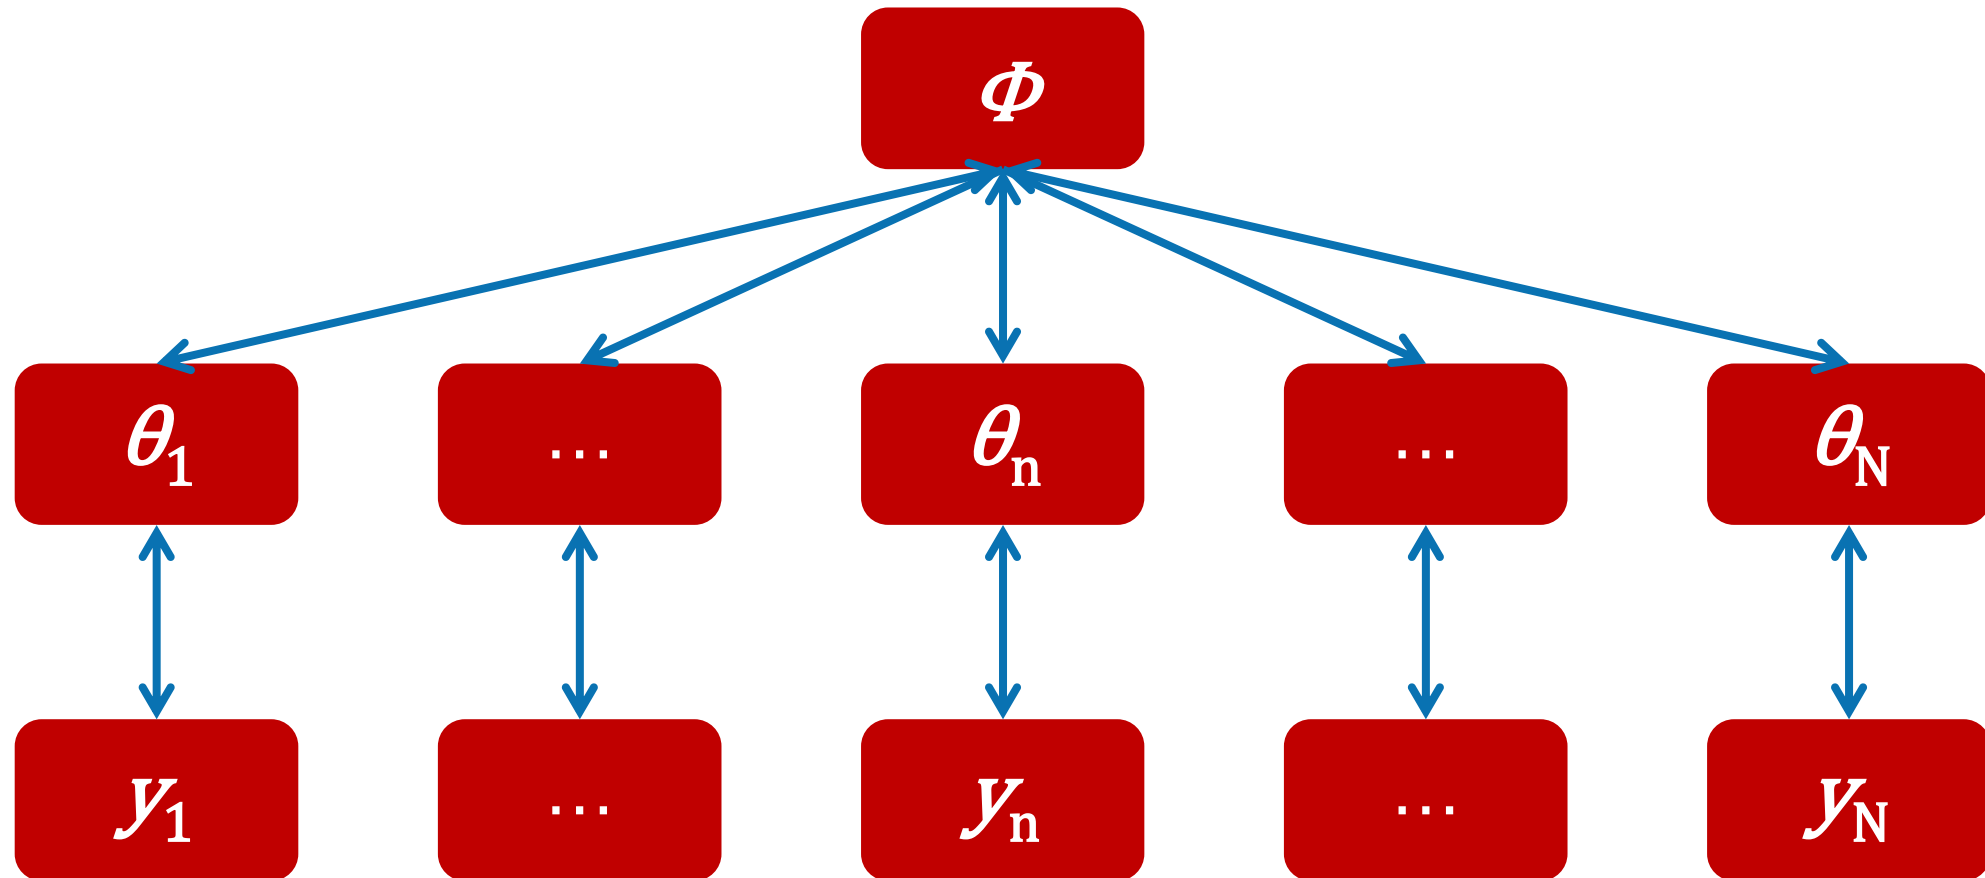


Hierarchical Structure

cognitive model

statistics

computing

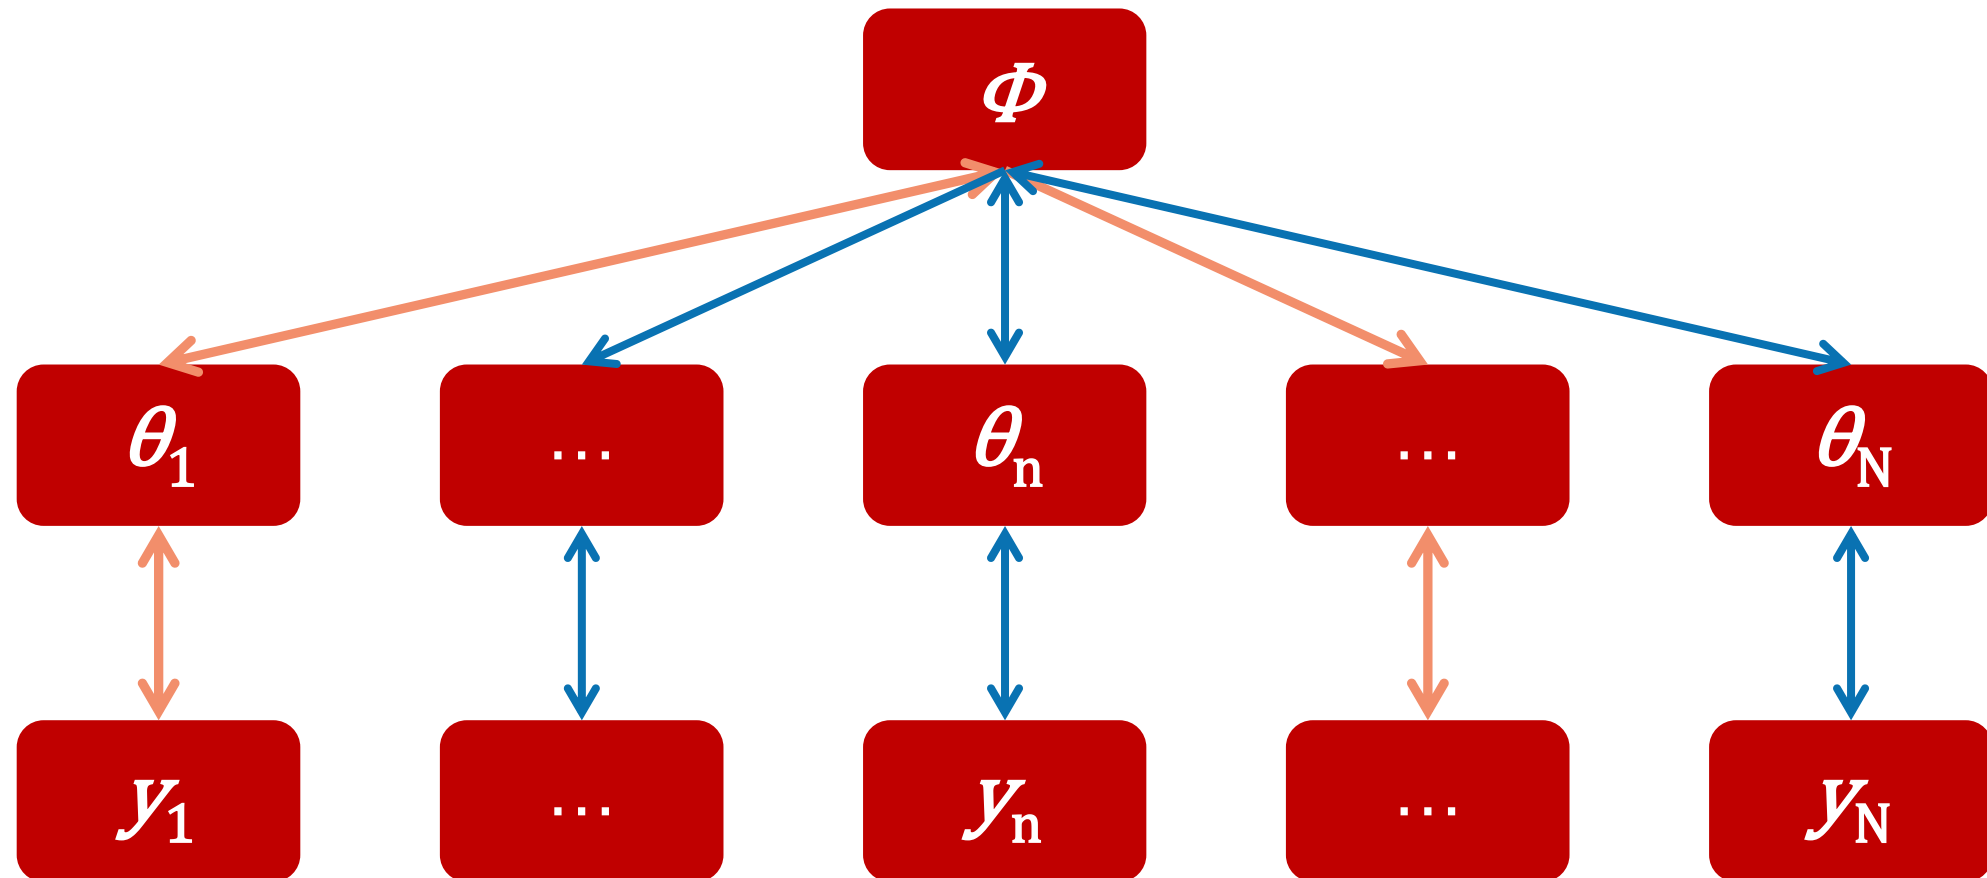


Hierarchical Structure

cognitive model

statistics

computing

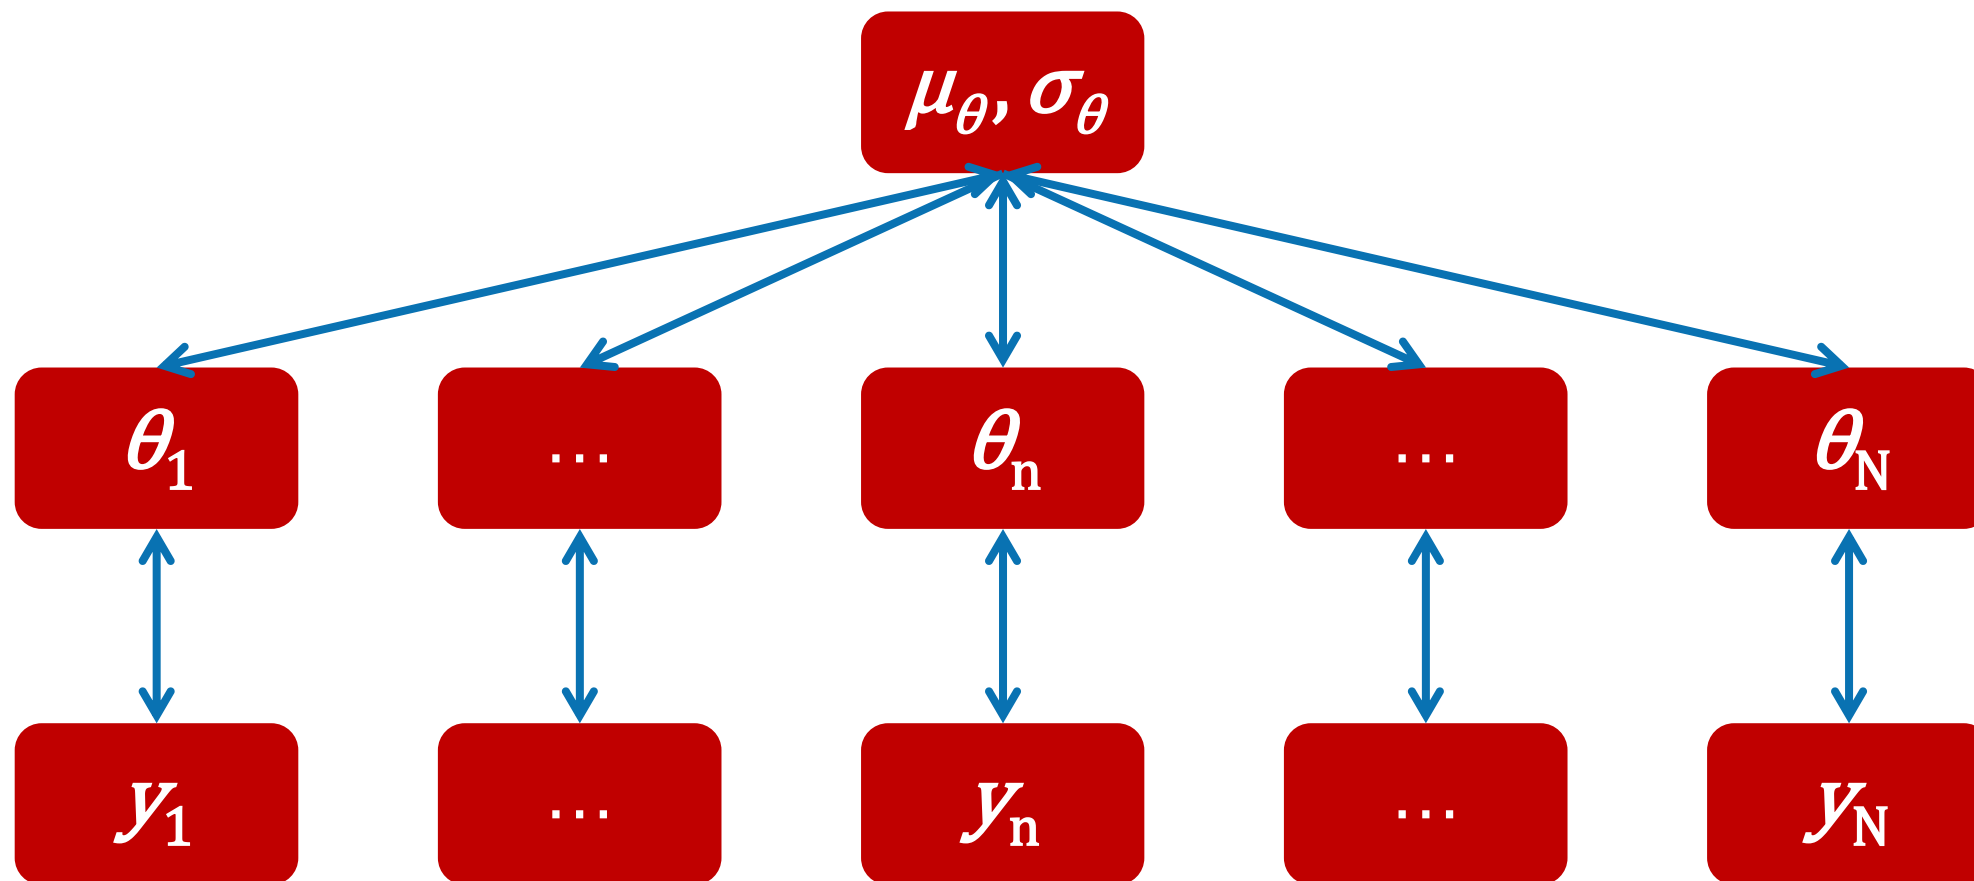


Hierarchical Structure

cognitive model

statistics

computing

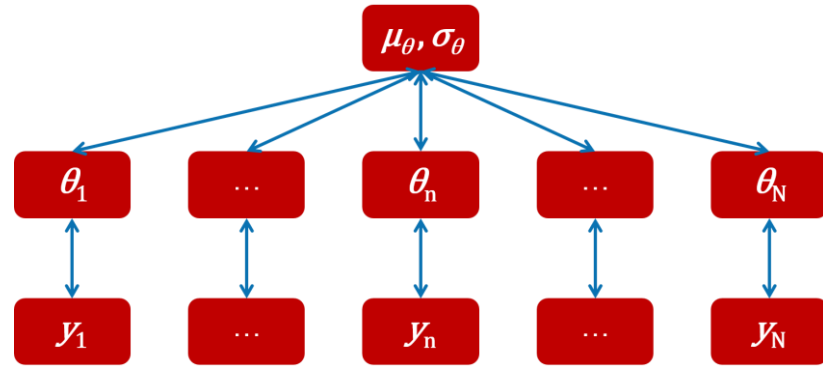


Hierarchical Structure

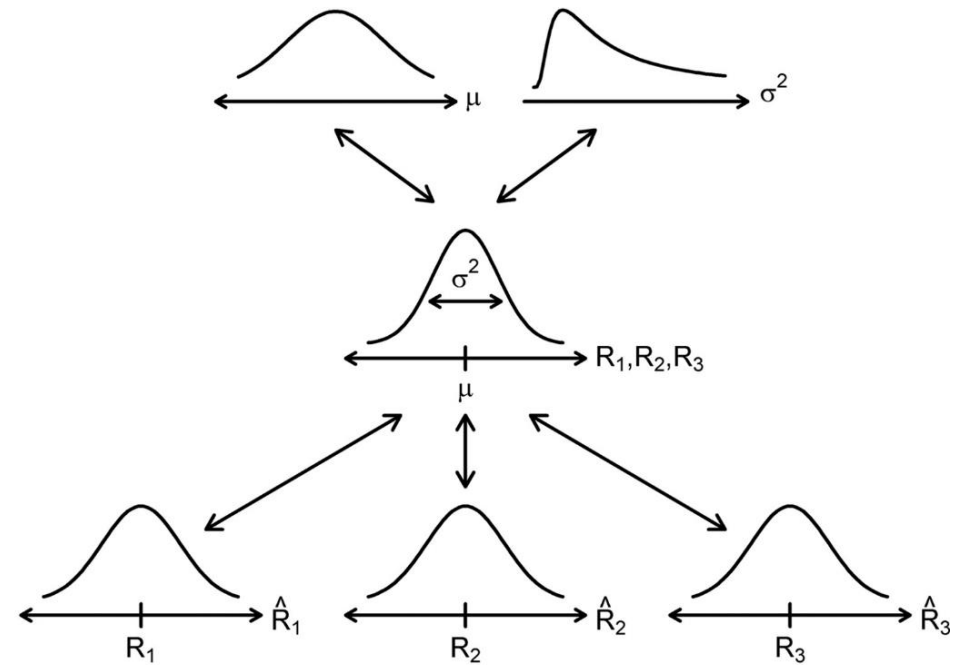
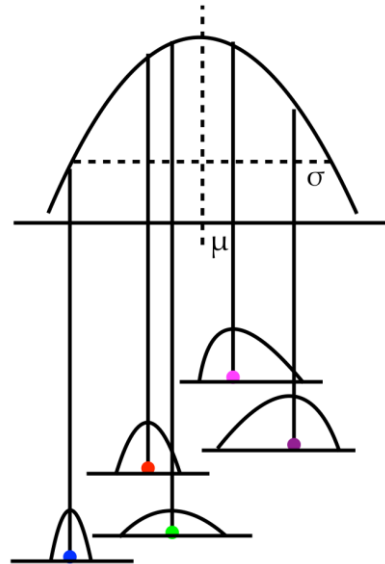
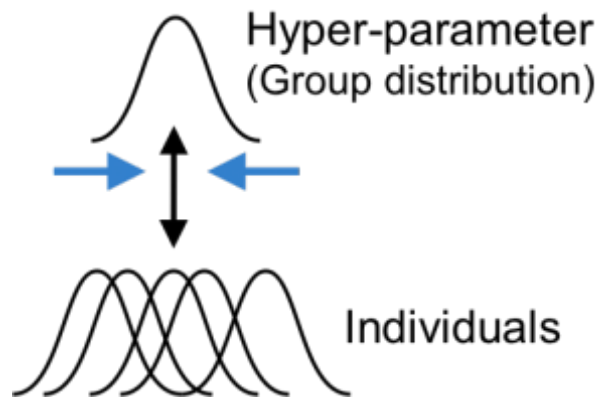
cognitive model

statistics

computing



$$P(\Theta, \Phi | D) = \frac{P(D | \Theta, \Phi) P(\Theta, \Phi)}{P(D)} \propto P(D | \Theta) P(\Theta | \Phi) P(\Phi)$$

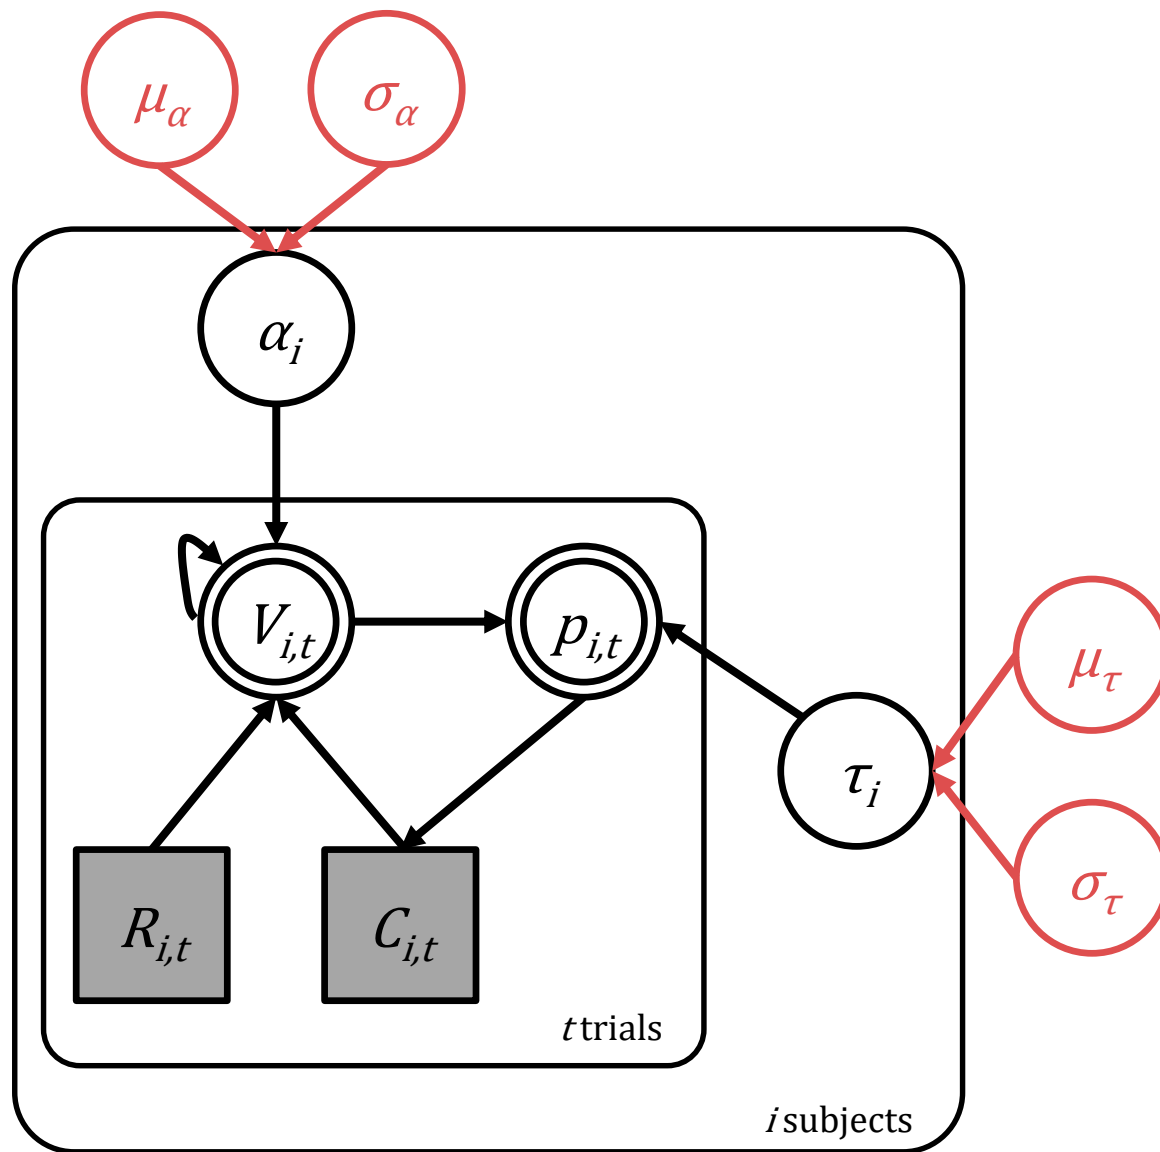


Hierarchical RL Model

cognitive model

statistics

computing

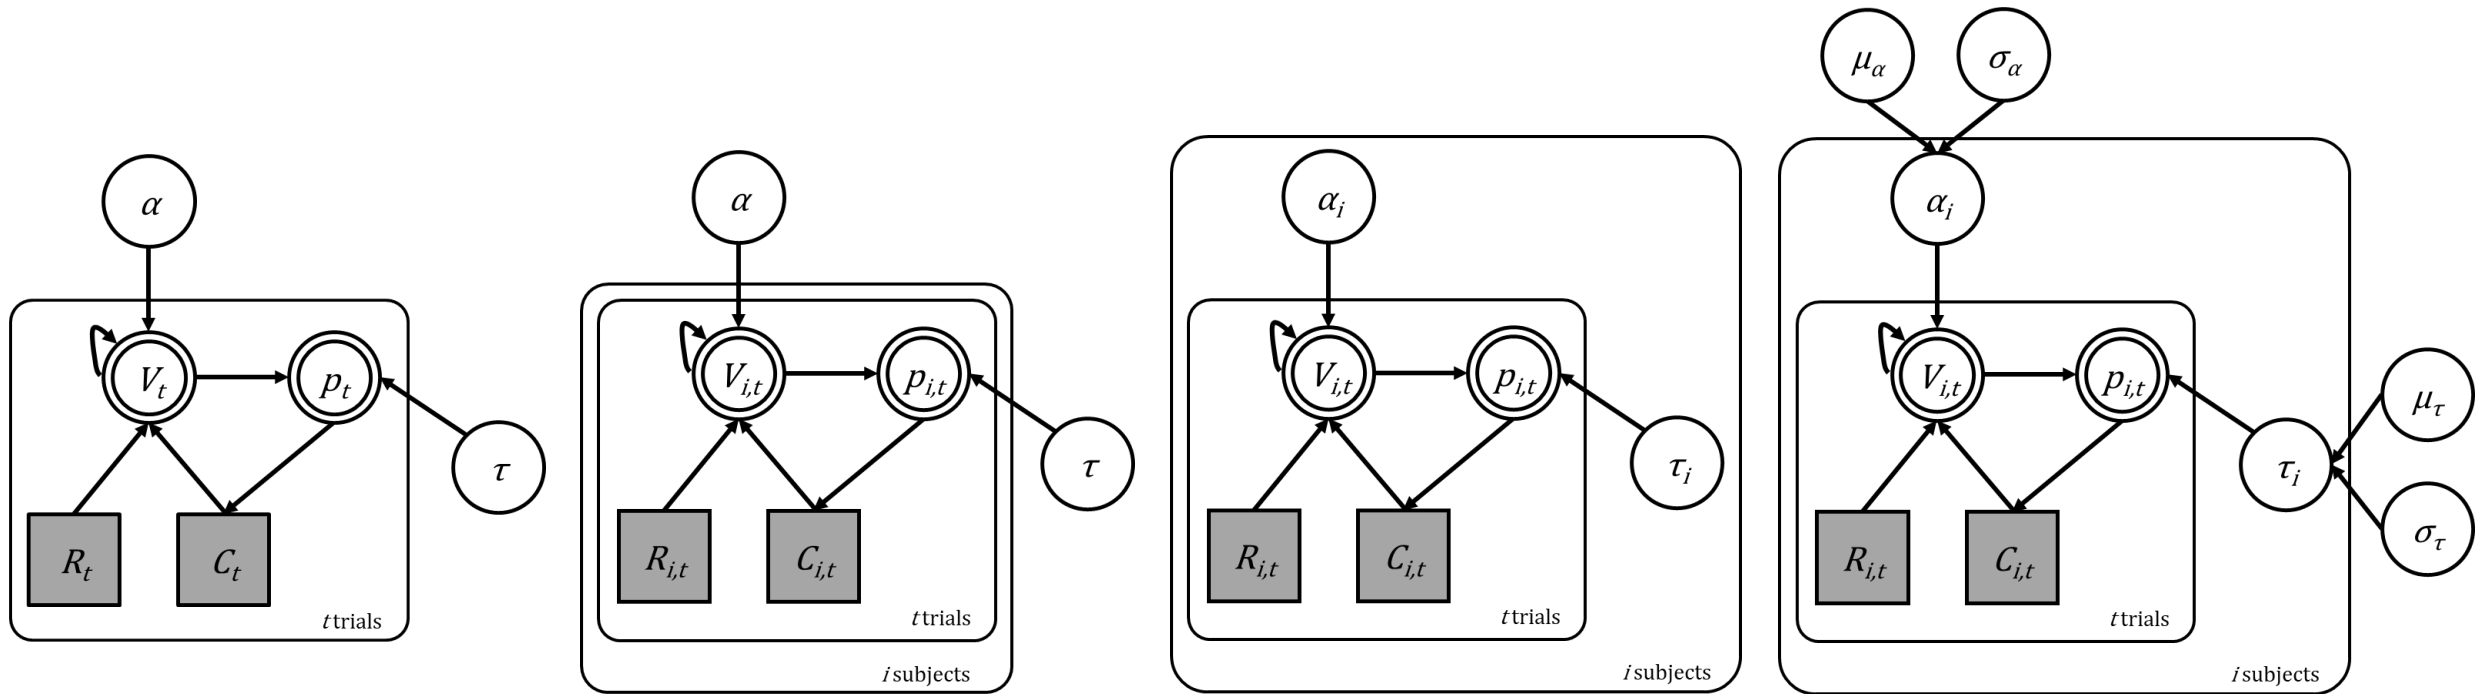


HOW DID WE GET HERE?

cognitive model

statistics

computing



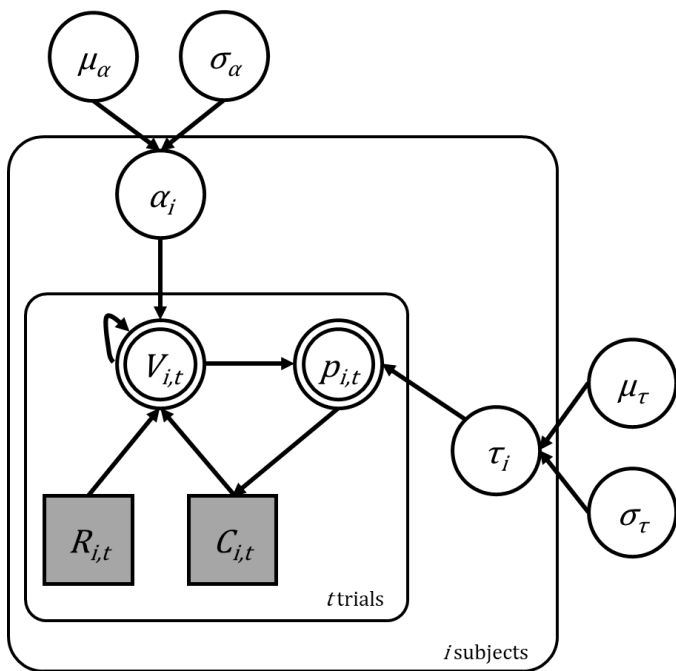
The cognitive model *per se* is the same!

Implementing Hierarchical RL Model

cognitive model

statistics

computing



$$\mu_{\alpha} \sim \text{Uniform}(0, 1)$$

$$\sigma_{\alpha} \sim \text{halfCauchy}(0, 1)$$

$$\mu_{\tau} \sim \text{Uniform}(0, 3)$$

$$\sigma_{\tau} \sim \text{halfCauchy}(0, 3)$$

$$\alpha_i \sim \text{Normal}(\mu_{\alpha}, \sigma_{\alpha}) \tau(0, 1)$$

$$\tau_i \sim \text{Normal}(\mu_{\tau}, \sigma_{\tau}) \tau(0, 3)$$

$$p_{i,t}(C = A) = \frac{1}{1 + e^{\tau_i(V_{i,t}(B) - V_{i,t}(A))}}$$

$$V_{i,t+1}^C = V_{i,t}^C + \alpha_i(R_{i,t} - V_{i,t}^C)$$

```
parameters {
  real<lower=0,upper=1> lr_mu;
  real<lower=0,upper=3> tau_mu;

  real<lower=0> lr_sd;
  real<lower=0> tau_sd;

  real<lower=0,upper=1> lr[nSubjects];
  real<lower=0,upper=3> tau[nSubjects];
}
```

```
model {
  lr_sd ~ cauchy(0, 1);
  tau_sd ~ cauchy(0, 3);
  lr ~ normal(lr_mu, lr_sd);
  tau ~ normal(tau_mu, tau_sd);
}
```

```
for (s in 1:nSubjects) {
  vector[2] v;
  real pe;
  v = initV;

  for (t in 1:nTrials) {
    choice[s,t] ~ categorical_logit( tau[s] * v );
    pe = reward[s,t] - v[choice[s,t]];
    v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;
  }
}
```

Exercise XI

cognitive model

statistics

computing

```
.../06.reinforcement_learning/_scripts/reinforcement_learning_multi_parm_main.R
```

TASK: (1) complete the model (TIP: individual ~ group)
(2) fit the hierarchical RL model

```
> source('_scripts/reinforcement_learning_multi_parm_main.R')  
  
> fit_rl3 <- run_rl_mp( modelType = 'hrch' )
```

In addition: Warning messages:

1: There were 97 divergent transitions after warmup. Increasing adapt_delta above 0.8 may help. See <http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup>

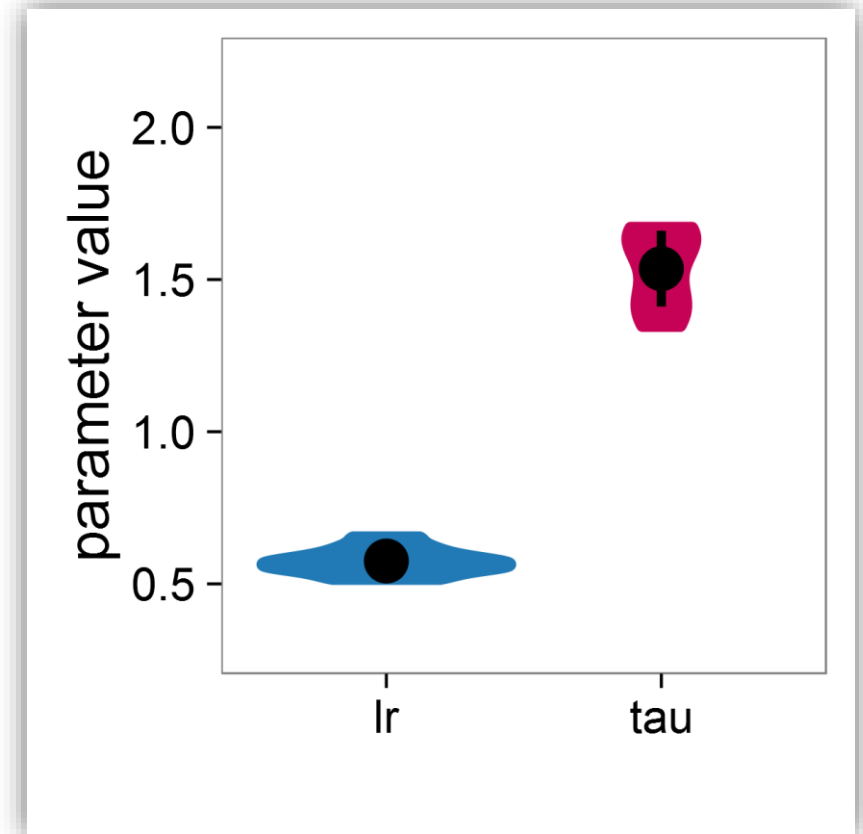
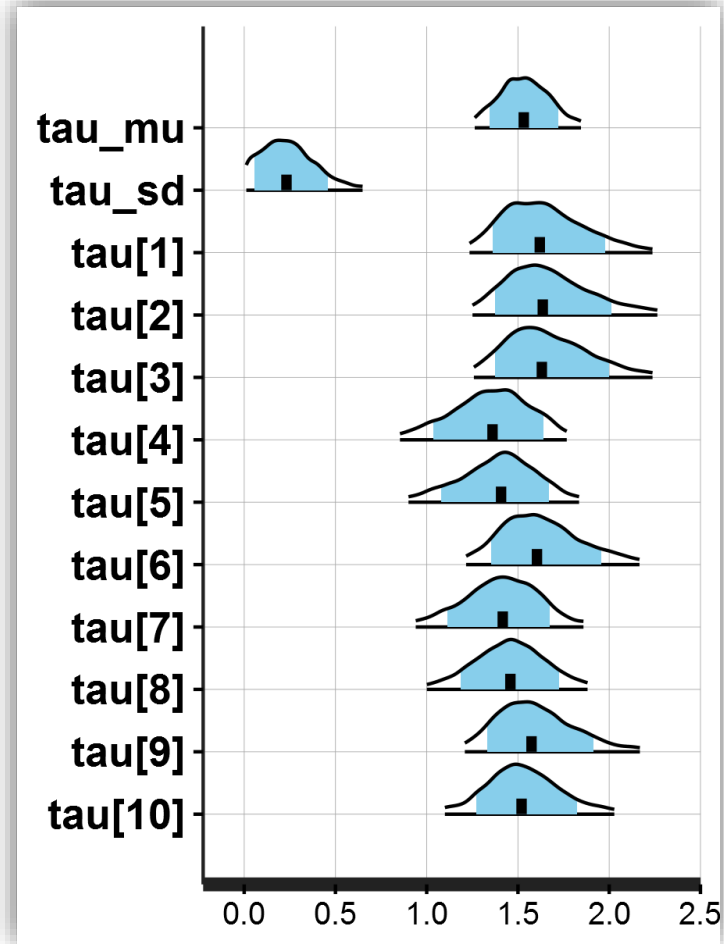
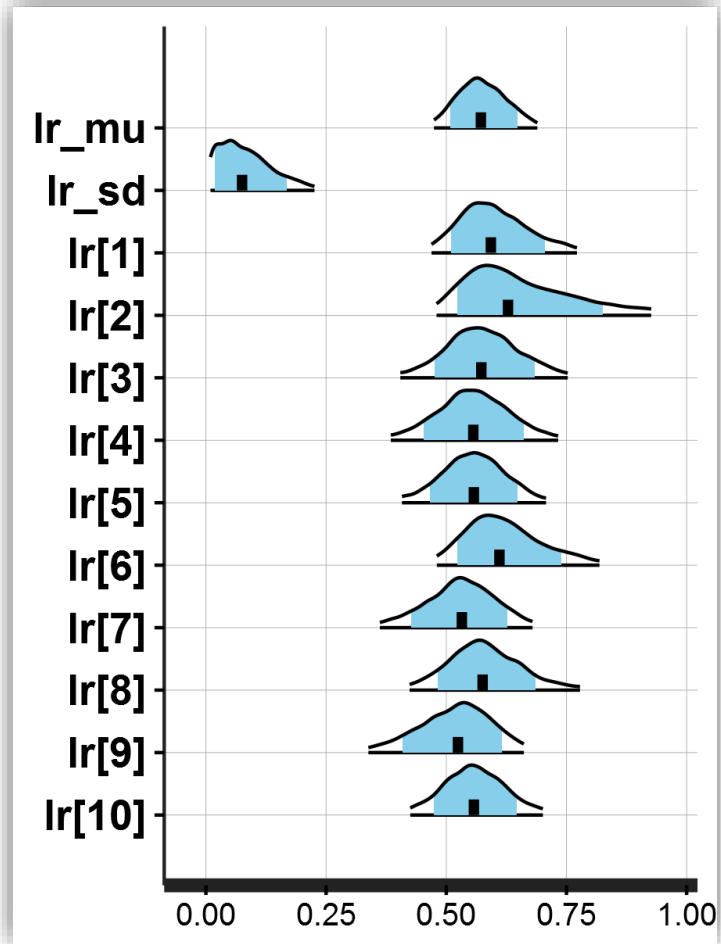
2: Examine the pairs() plot to diagnose sampling problems

Hierarchical Fitting*

cognitive model

statistics

computing



*: adapt_delta=0.999, max_treedepth=100

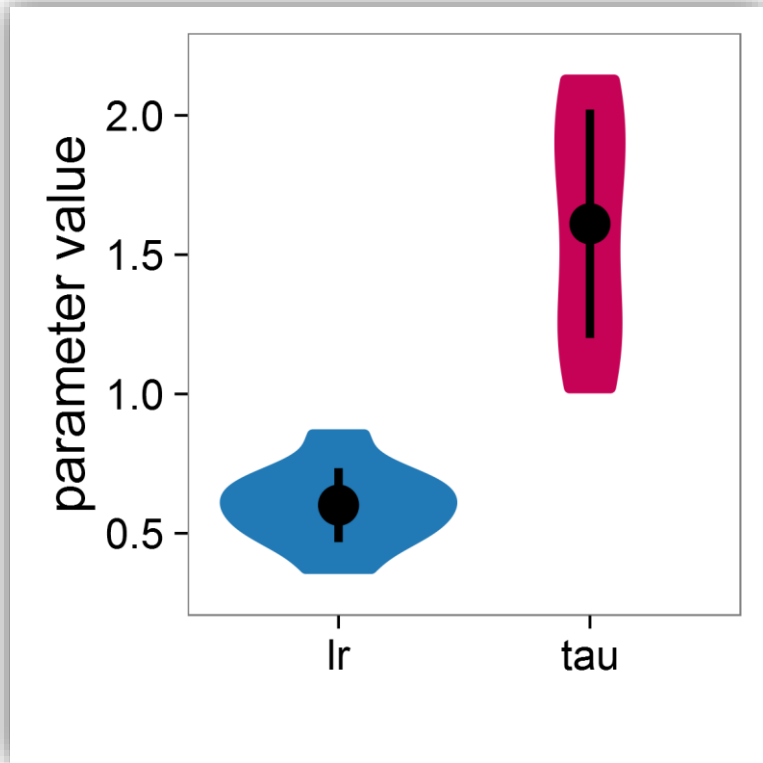
Comparing with True Parameters

cognitive model

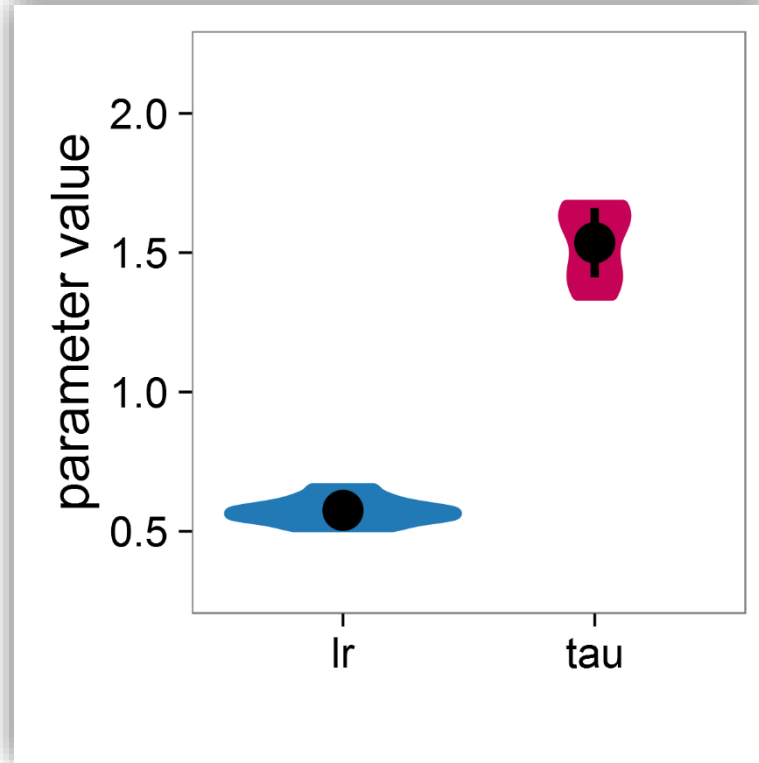
statistics

computing

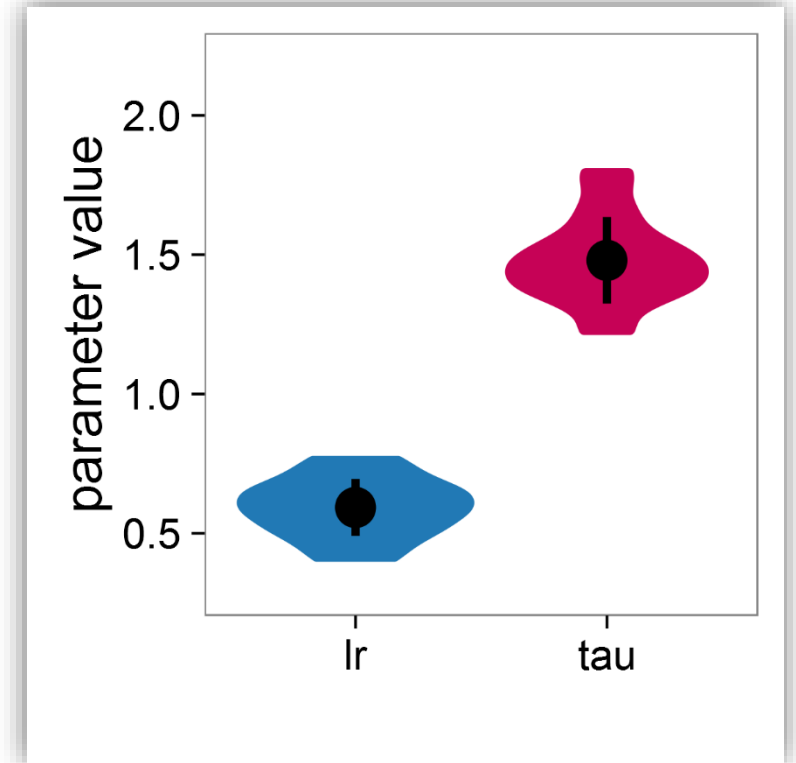
Posterior Means (indv)



Posterior Means (hrch)*



True Parameters



*: adapt_delta=0.999, max_treedepth=100

Group-level Parameters

cognitive model

statistics

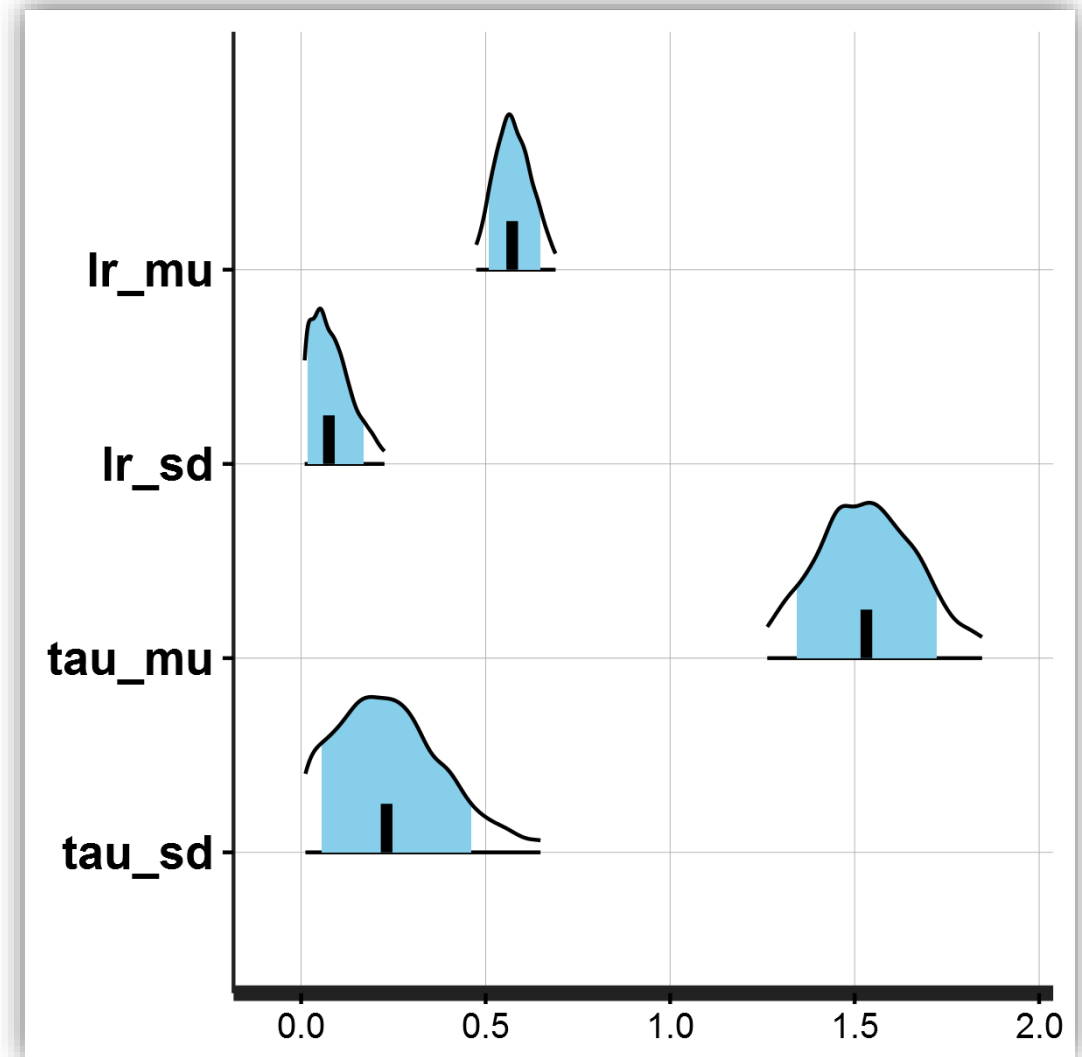
computing

True group parameters

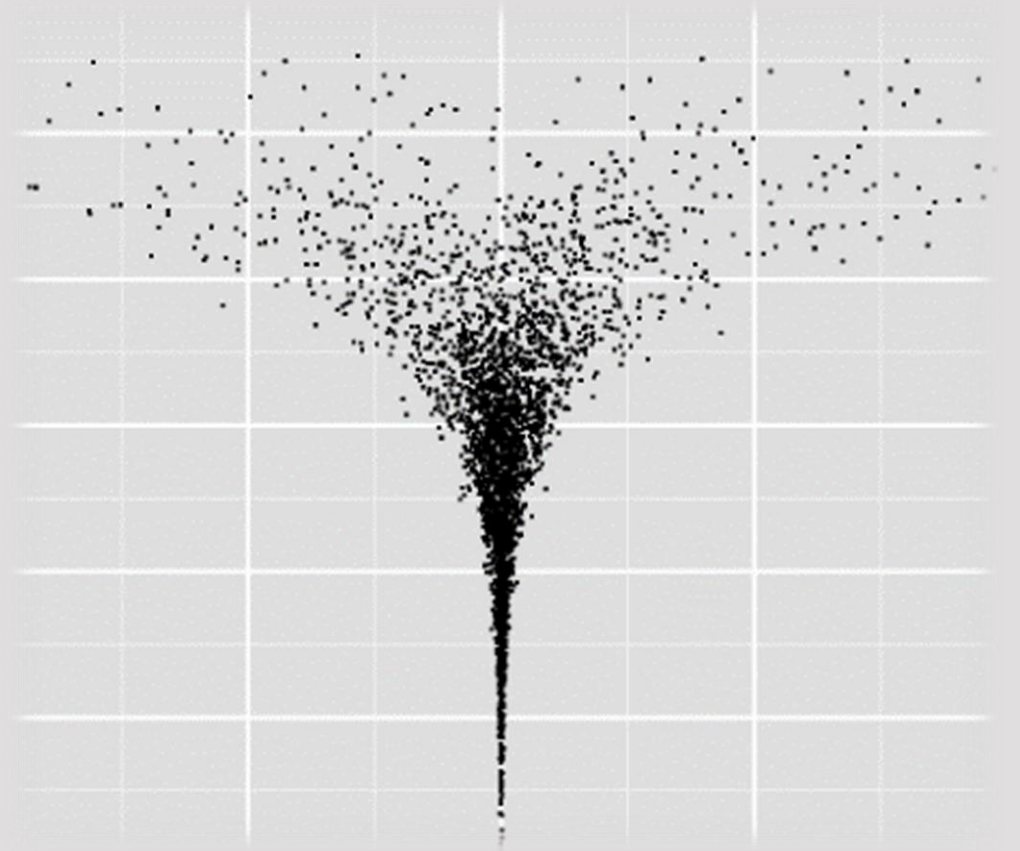
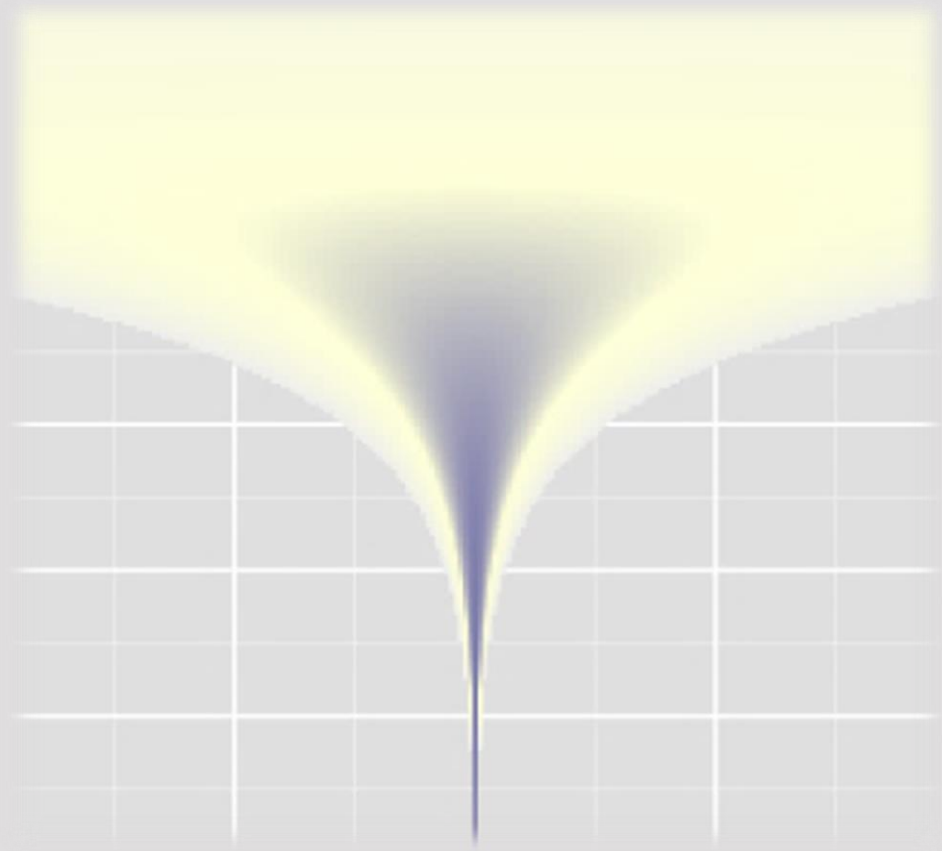
```
lr = rnorm(10, mean=0.6, sd=0.12)
tau = rnorm(10, mean=1.5, sd=0.2)
```

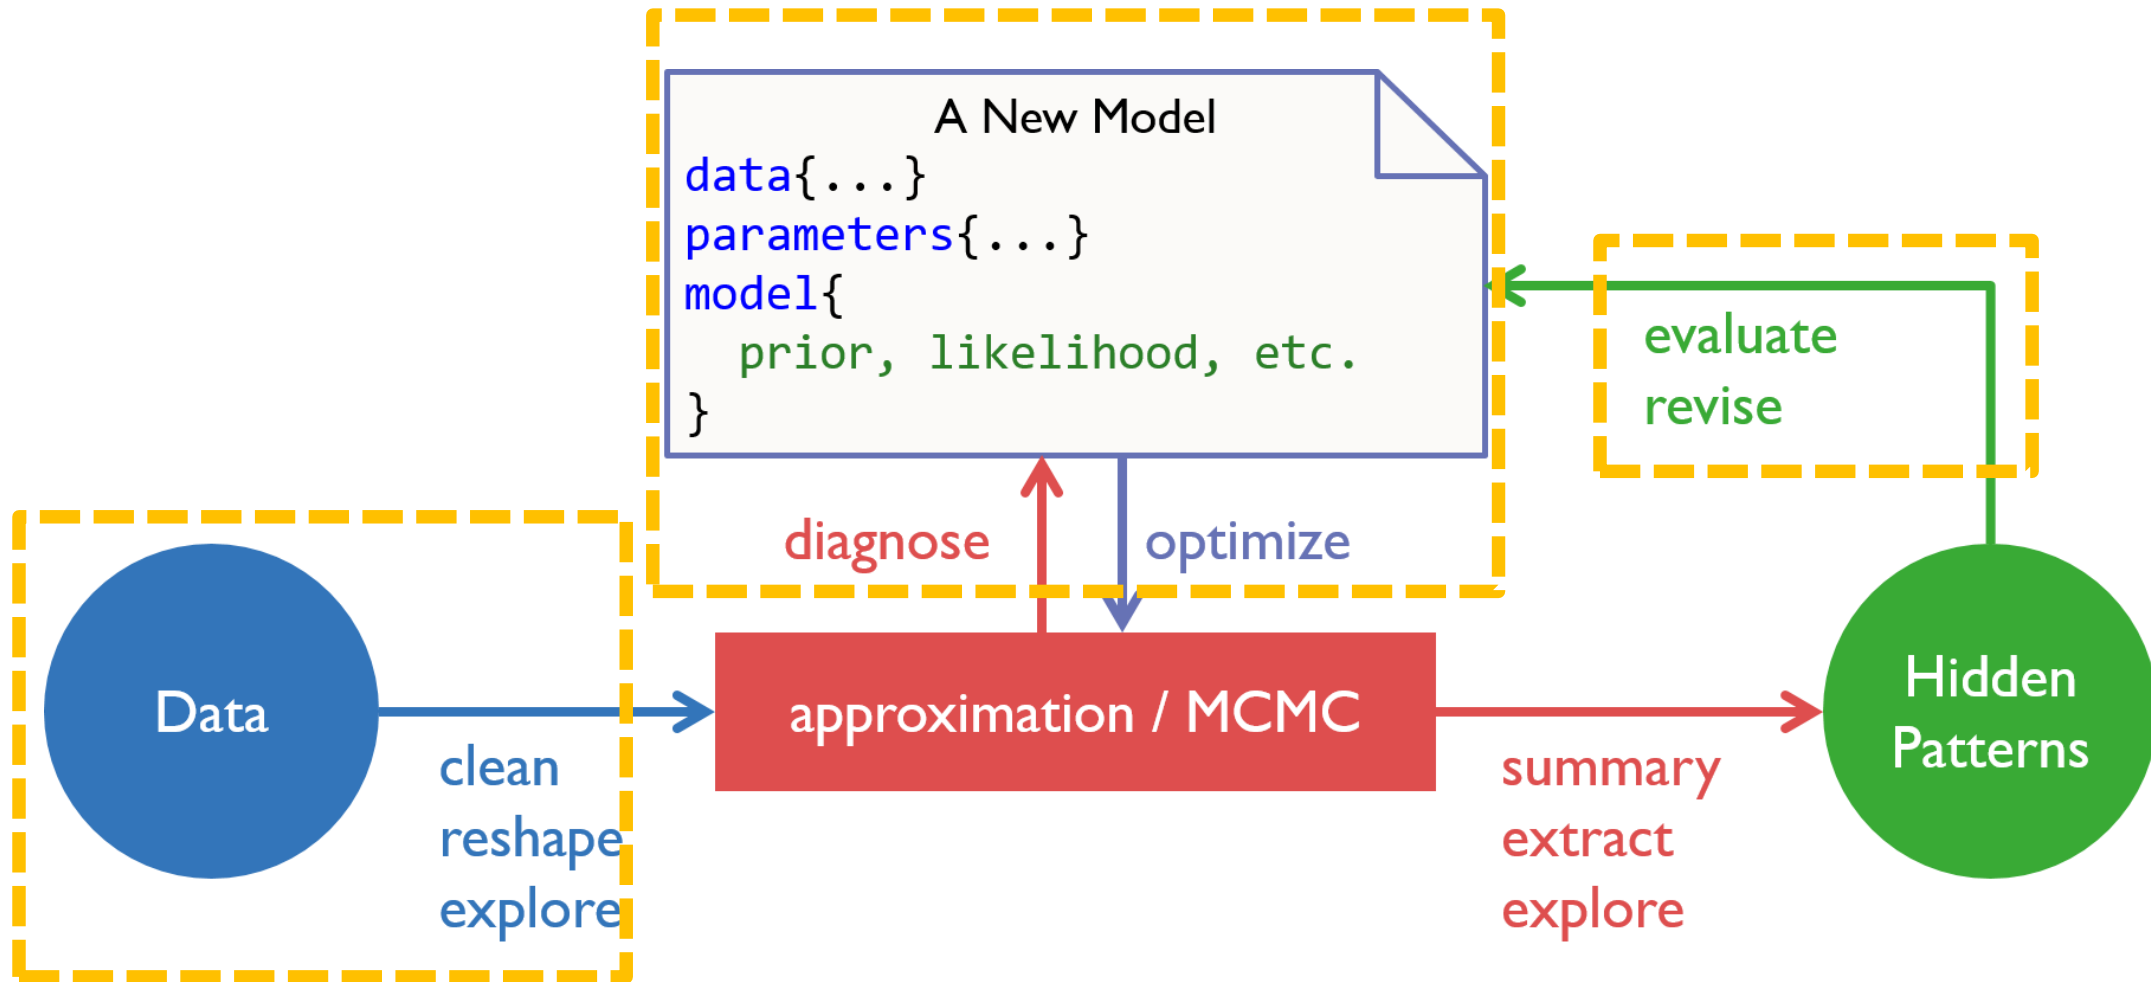
Estimated group parameters

	mean	2.5%	25%	50%	75%	97.5%
lr_mu	0.58	0.47	0.54	0.57	0.61	0.69
lr_sd	0.09	0.01	0.04	0.08	0.12	0.23
tau_mu	1.54	1.26	1.43	1.53	1.63	1.85
tau_sd	0.25	0.01	0.13	0.23	0.34	0.65



OPTIMIZING STAN CODES







Optimizing Stan Code

cognitive model

statistics

computing

Preprocess data

run as many calculations as you can outside Stan

Specify a proper model

follow literature, supervision, experience, etc.

Vectorizing

vectorize Stan code whenever you can

Reparameterizing

reparameterize target parameter to simple distributions

Preprocess Data

cognitive model

statistics

computing

$$\overline{\text{height}} = \alpha + \beta_1 * \text{weight} + \beta_2 * \text{weight}^2$$

```
d$weight_sq <- d$weight^2
```

```
data {  
  int<lower=0> N;  
  vector<lower=0>[N] height;  
  vector<lower=0>[N] weight;  
  vector<lower=0>[N] weight_sq;  
}
```

Specify a Proper Model

cognitive model

statistics

computing

- Visualize your data
- Follow Literatures
- Start from simple and then build complexities
- Simulate data and run model recovery

A New Model

```
data{...}  
parameters{...}  
model{  
  prior, likelihood, etc.  
}
```

Vectorization

cognitive model

statistics

computing

```
model {  
  for (n in 1:N) {  
    flip[n] ~ bernoulli(theta);  
  }  
}
```



```
model {  
  flip ~ bernoulli(theta);  
}
```

```
parameters {  
  ...  
  real<lower=0,upper=1> lr[nSubjects];  
  real<lower=0,upper=3> tau[nSubjects];  
}  
  
model {  
  ...  
  lr ~ normal(lr_mu, lr_sd) ;  
  tau ~ normal(tau_mu, tau_sd) ;  
  ...  
}
```

```
model {  
  vector[N] mu;  
  for (i in 1:N) {  
    mu[i] = alpha + beta * weight[i];  
    height[i] ~ normal(mu[i], sigma)  
  }  
}
```



```
model {  
  vector[N] mu;  
  mu = alpha + beta * weight;  
  height ~ normal(mu, sigma);  
}
```



```
model {  
  height ~ normal(alpha + beta * weight, sigma);  
}
```

Reparameterization

Neal's Funnel

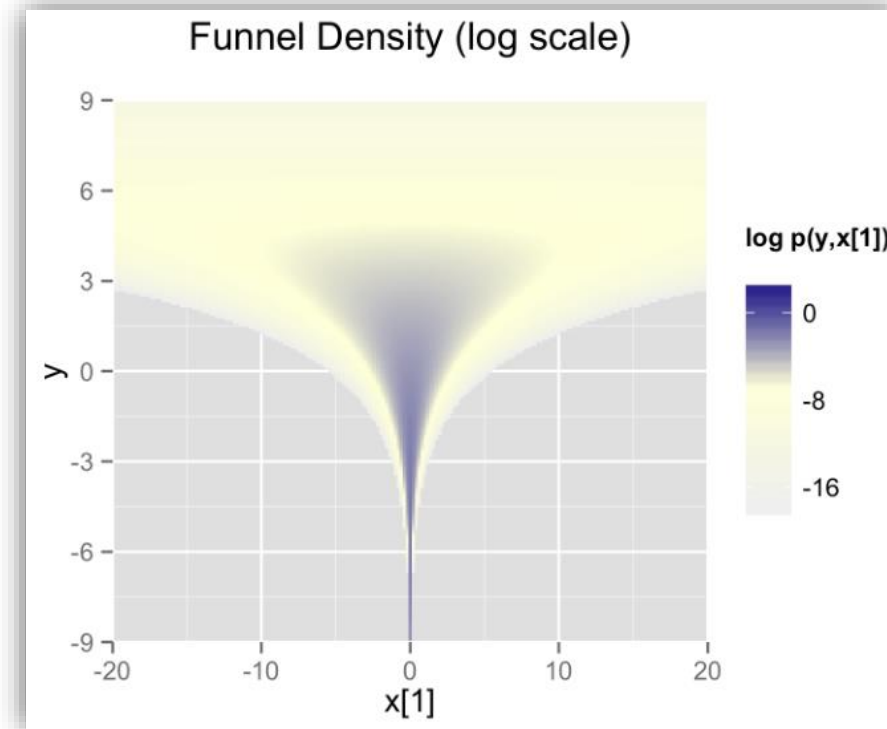
cognitive model

statistics

computing

$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```



Non-centered Reparameterization*

cognitive model

statistics

computing

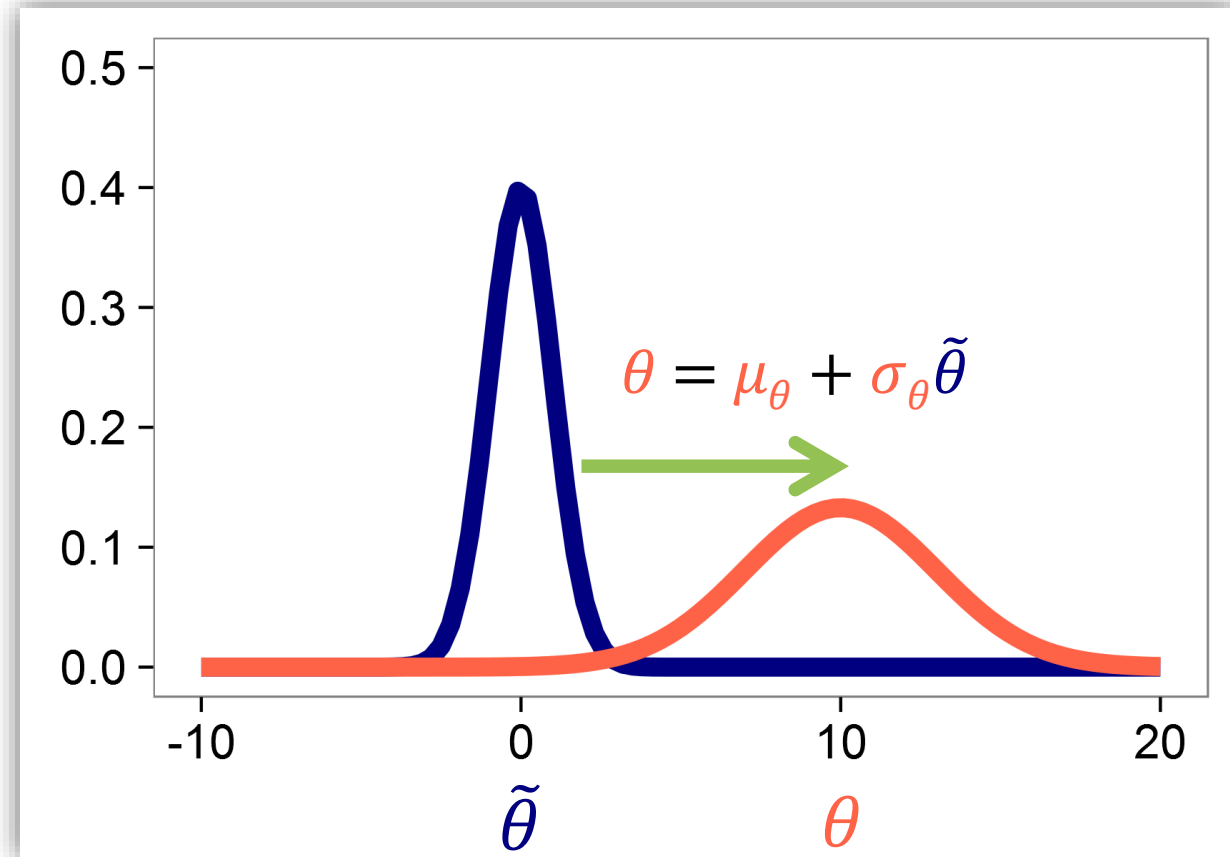
$$\theta \sim \text{Normal}(\mu_\theta, \sigma_\theta)$$



$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \mu_\theta + \sigma_\theta \tilde{\theta}$$

Stan likes **simple**
distributions!



Reparameterization

Neal's Funnel

cognitive model

statistics

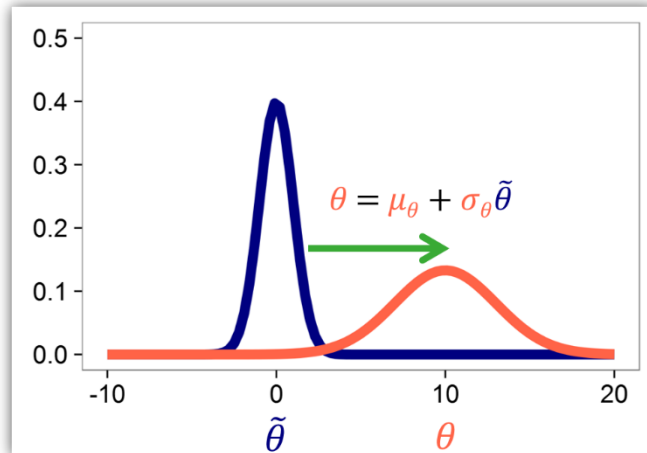
computing

$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```



```
parameters {  
  real y_raw;  
  vector[9] x_raw;  
}  
transformed parameters {  
  real y;  
  vector[9] x;  
  
  y = 3.0 * y_raw;  
  x = exp(y/2) * x_raw;  
}  
model {  
  y_raw ~ normal(0, 1);  
  x_raw ~ normal(0, 1);  
}
```



Stan Sampling Parameters

cognitive model

statistics

computing

parameter	description	constraint	default
iterations	number of MCMC samples (per chain)	int, > 0	2000
delta: δ	target Metropolis acceptance rate	$\delta \in [0, 1]$	0.80
stepsize: ε	initial HMC step size	real, $\varepsilon > 0$	2.0
max_treedepth: L	maximum HMC steps per iteration	int, $L > 0$	10

Typical adjustments

- Increase iterations
- Increase delta
- Decrease stepsize
- Might have to increase max_treedepth

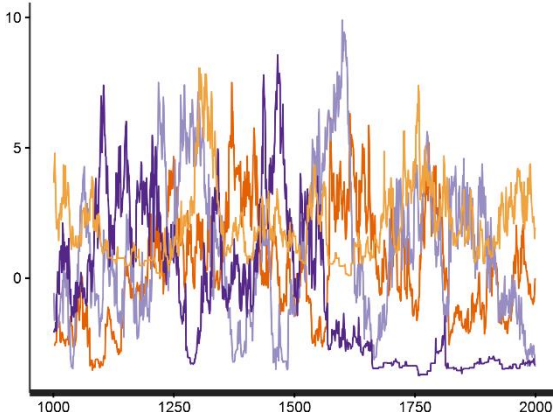
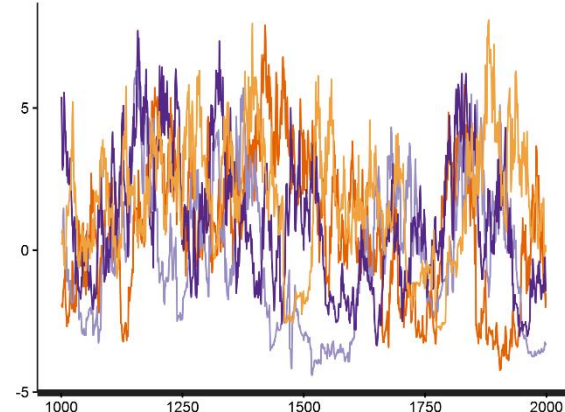
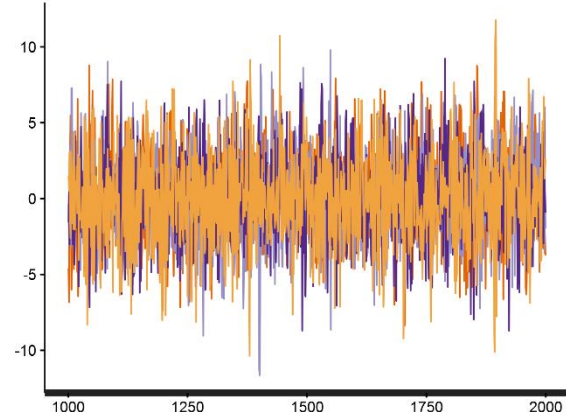
```
funnel_fit2 <- stan("_scripts/funnel.stan",  
  iter = 4000,  
  control = list(adapt_delta = 0.999,  
    stepsize = 1.0,  
    max_treedepth = 20))
```

Neal's Funnel: Comparing Performance

cognitive model

statistics

computing

	direct model	adjusted direct model	reparameterized model
Rhat (y)	1.22	1.1	1.0
n_eff (y)	18	42	3886
runtime*	48.50 sec	50.76 sec	50.12 sec
n_eff (y) / runtime	0.37 / sec	0.82 / sec	77.53 / sec
n_divergent	53	0	0
traceplot (y)			

*: 2 cores in parallel, including compiling time

How about **Bounded** Parameters?

$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \mu_{\theta} + \sigma_{\theta} \tilde{\theta}$$

$$\theta \in (-\infty, +\infty)$$



$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta})$$

$$\theta \in [0, 1]$$

constraint	reparameterization
$\theta \in (-\infty, +\infty)$	$\theta = \mu_{\theta} + \sigma_{\theta} \tilde{\theta}$
$\theta \in [0, N]$	$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \times N$
$\theta \in [M, N]$	$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \times (N-M) + M$
$\theta \in (0, +\infty)$	$\theta = \exp(\mu_{\theta} + \sigma_{\theta} \tilde{\theta})$

* Probit⁻¹: Normal cumulative distribution function (normcdf)

Apply to Our Hierarchical RL Model

cognitive model

statistics

computing

```
parameters {  
  real<lower=0,upper=1> lr_mu;  
  real<lower=0,upper=3> tau_mu;  
  
  real<lower=0> lr_sd;  
  real<lower=0> tau_sd;  
  
  real<lower=0,upper=1> lr[nSubjects];  
  real<lower=0,upper=3> tau[nSubjects];  
}
```



```
parameters {  
  # group-level parameters  
  real lr_mu_raw;  
  real tau_mu_raw;  
  real<lower=0> lr_sd_raw;  
  real<lower=0> tau_sd_raw;  
  
  # subject-level raw parameters  
  vector[nSubjects] lr_raw;  
  vector[nSubjects] tau_raw;  
}  
  
transformed parameters {  
  vector<lower=0,upper=1>[nSubjects] lr;  
  vector<lower=0,upper=3>[nSubjects] tau;  
  
  for (s in 1:nSubjects) {  
    lr[s] = Phi_approx( lr_mu_raw + lr_sd_raw * lr_raw[s] );  
    tau[s] = Phi_approx( tau_mu_raw + tau_sd_raw * tau_raw[s] ) * 3;  
  }  
}
```

Apply to Our Hierarchical RL Model

cognitive model

statistics

computing

```
model {  
  lr_sd ~ cauchy(0,1);  
  tau_sd ~ cauchy(0,3);  
  lr ~ normal(lr_mu, lr_sd) ;  
  tau ~ normal(tau_mu, tau_sd) ;  
  
  for (s in 1:nSubjects) {  
    vector[2] v;  
    real pe;  
    v = initV;  
  
    for (t in 1:nTrials) {  
      choice[s,t] ~ categorical_logit( tau[s] * v );  
      pe = reward[s,t] - v[choice[s,t]];  
      v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;  
    }  
  }  
}
```



```
model {  
  lr_mu_raw ~ normal(0,1);  
  tau_mu_raw ~ normal(0,1);  
  lr_sd_raw ~ cauchy(0,3);  
  tau_sd_raw ~ cauchy(0,3);  
  
  lr_raw ~ normal(0,1);  
  tau_raw ~ normal(0,1);  
  
  for (s in 1:nSubjects) {  
    ...  
  }  
  
  generated quantities {  
    real<lower=0,upper=1> lr_mu;  
    real<lower=0,upper=3> tau_mu;  
  
    lr_mu = Phi_approx(lr_mu_raw);  
    tau_mu = Phi_approx(tau_mu_raw) * 3;  
  }  
}
```

Exercise XII

cognitive model

statistics

computing

```
.../07.optm_rl/_scripts/reinforcement_learning_hrch_main.R
```

TASK: (1) Complete the Matt Trick
(2) fit the optimized hierarchical RL model

```
> source('_scripts/reinforcement_learning_hrch_main.R')  
> fit_rl4 <- run_rl_mp2(optimized = TRUE)
```

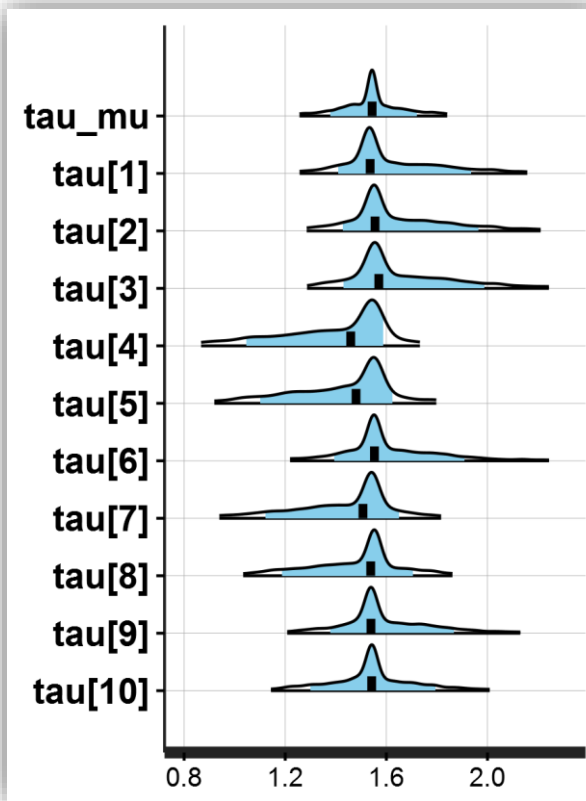
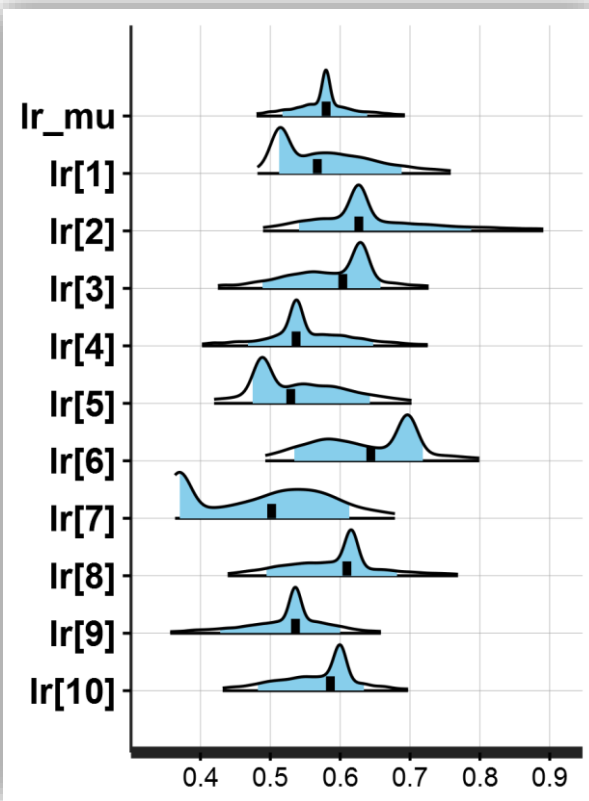
Hierarchical Fitting – Optimized

cognitive model

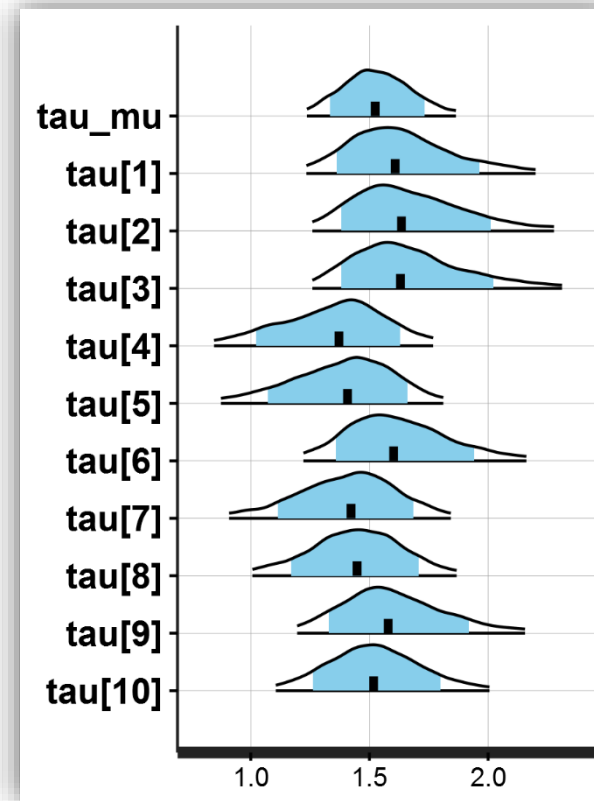
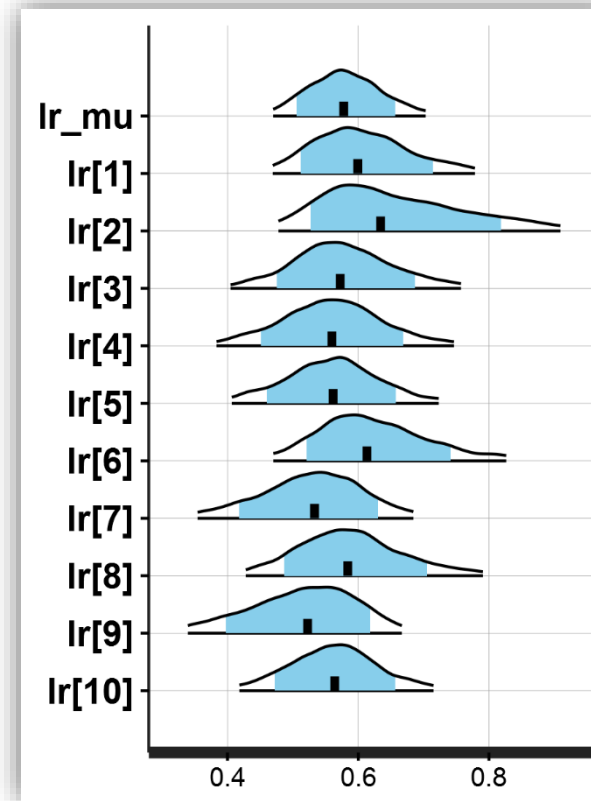
statistics

computing

Posterior Means (hrch)



Posterior Means (hrch + optm)



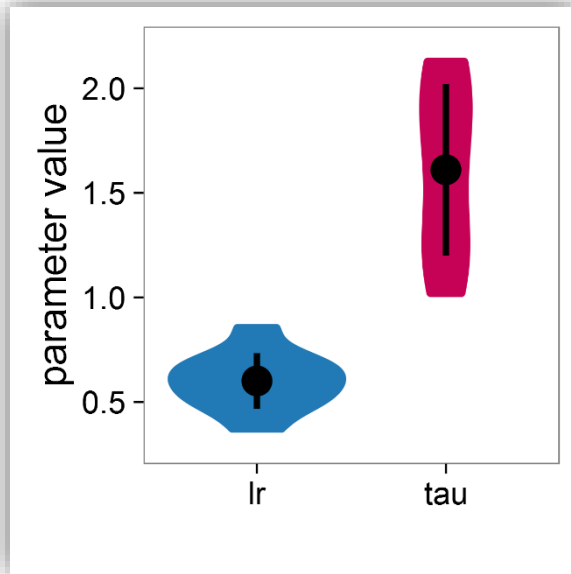
Comparing with True Parameters

cognitive model

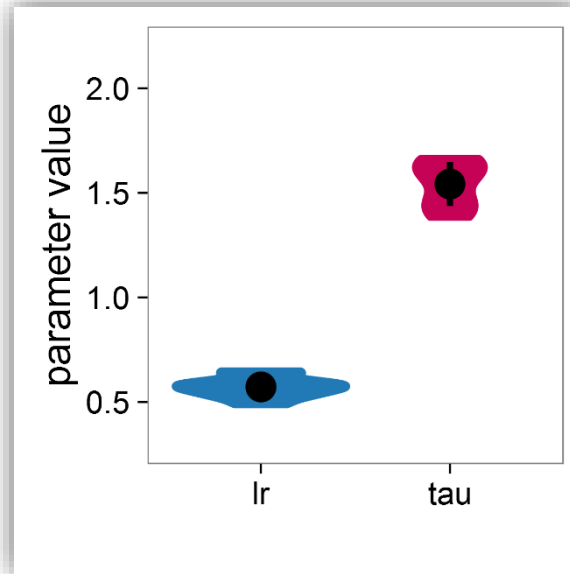
statistics

computing

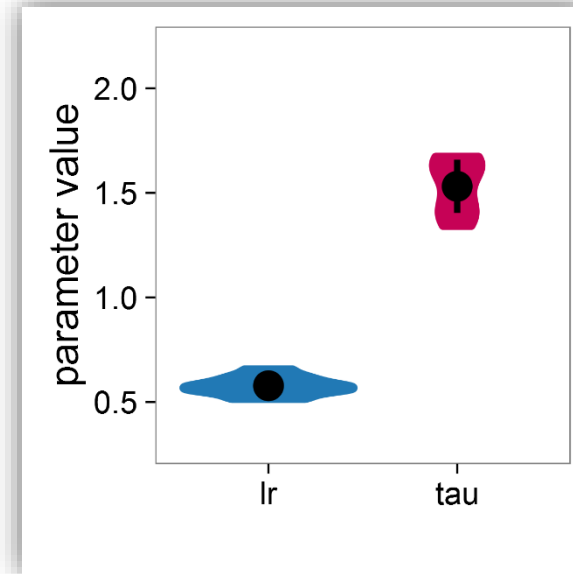
Posterior Means (indy)



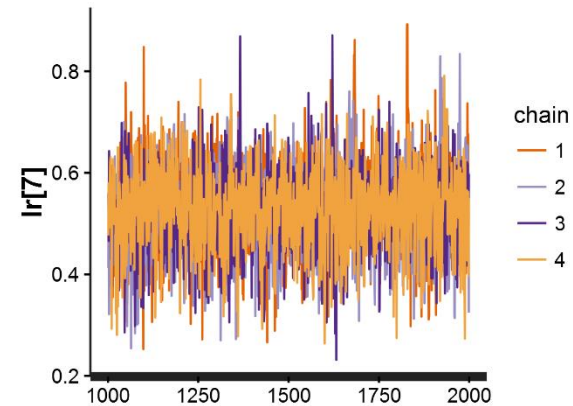
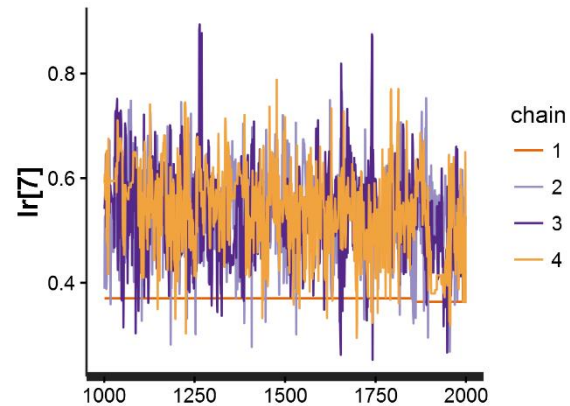
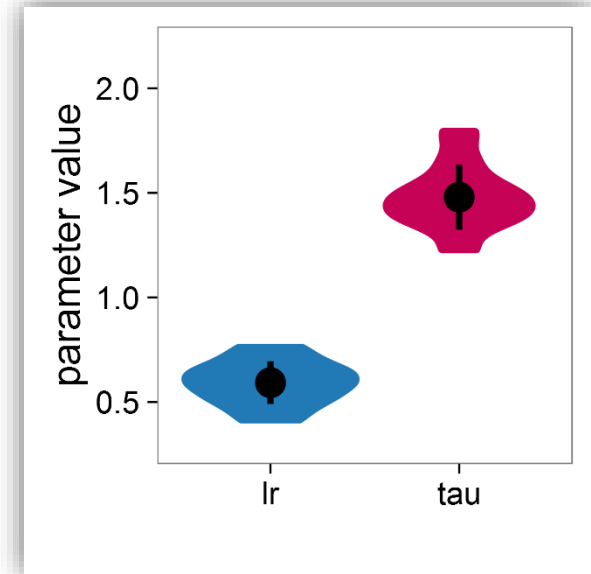
Posterior Means (hrch)



Posterior Means (hrch+optm)



True Parameters

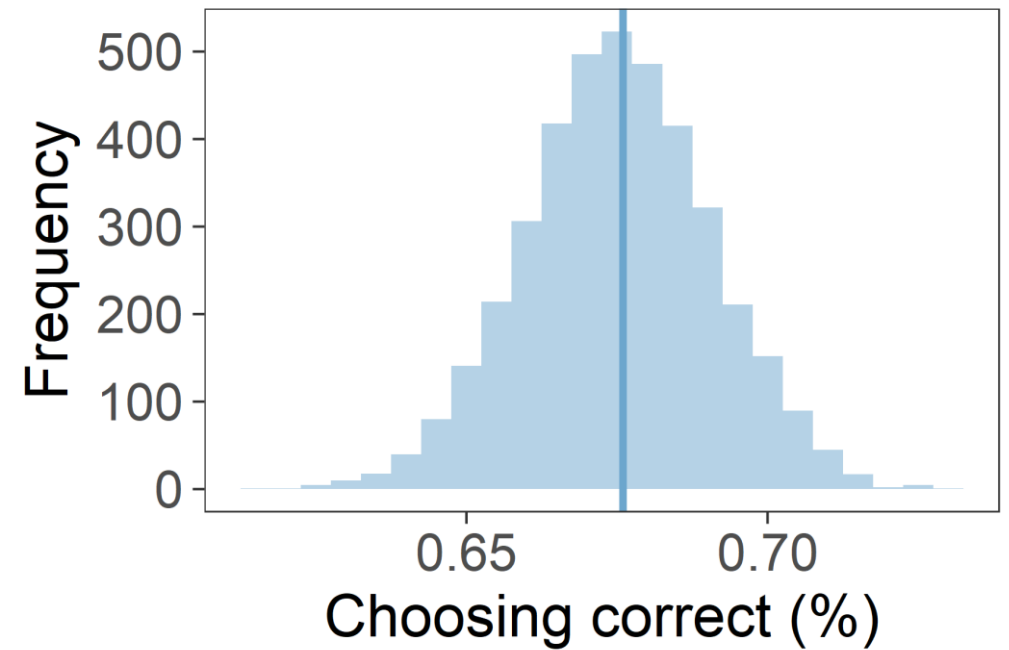
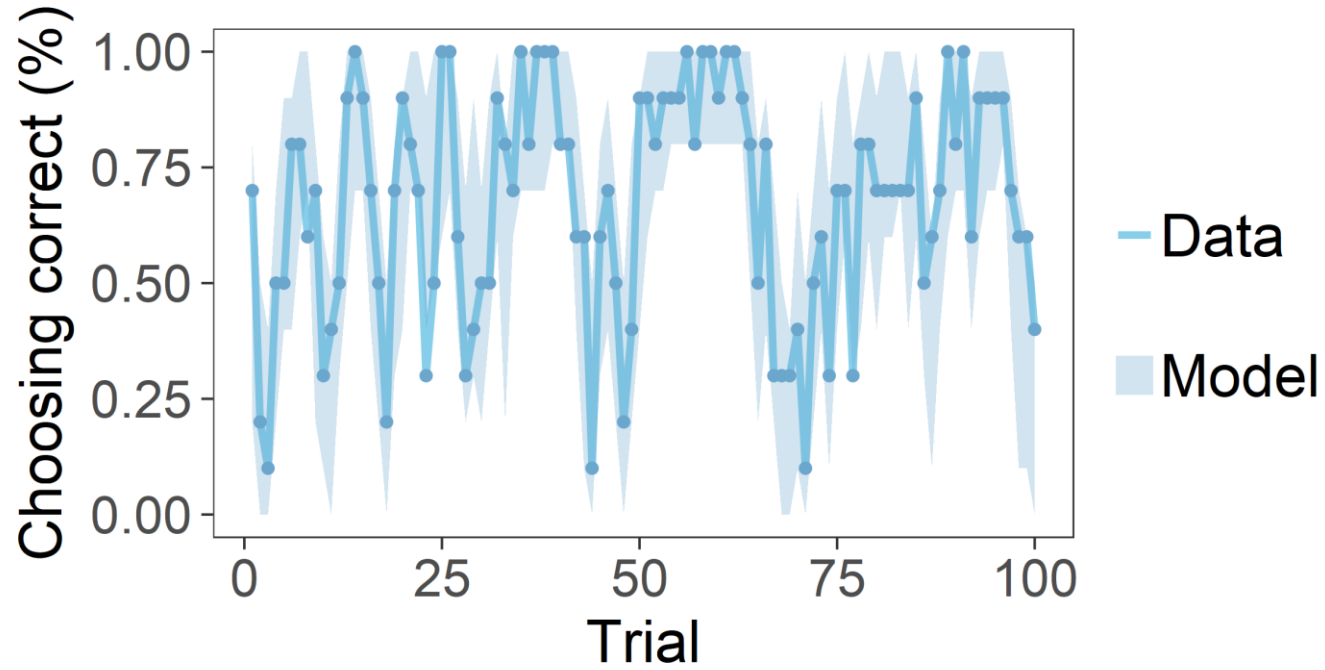


Posterior Predictive Check

cognitive model

statistics

computing



ANY
QUESTIONS
?

Happy Computing!